

Introduction to Matrix Factorization


David Rosenberg

New York University

December 26, 2016

The Matrix Denoising (or Smoothing) Problem

- Consider an image as a matrix of gray-scale values:



- Left: **Original**; Middle: **Adds Gaussian noise**; Right: **Denoised**

Taken from Marc'Aurelio Ranzato's page <http://www.cs.nyu.edu/~ranzato/research/projects.html>.

The Matrix Denoising (or Smoothing) Problem

- M is the **original matrix**
- W is the **noise** added to Z
- We observe $Z = M + W$.
- Problem:
 - Given Z , produce some estimate for M .
- How?

General Approach

- Observe Z (the noisy M matrix).
- We want a matrix \hat{Z} such that

$$\sum_{i=1}^m \sum_{j=1}^n (\hat{z}_{ij} - z_{ij})^2 \text{ is small.}$$

- But not too small!
 - Don't want to be fitting the noise (**overfitting**)
- Need some way to constrain \hat{Z} .
- Different constraints give different algorithms.

Some Constraints for \hat{M} (Quick Look)

- 1 $\|\hat{Z}\|_{\ell_1} \leq c$ (Sparse matrix approximation)
- 2 $\text{rank}(\hat{Z}) \leq c$. (SVD methods)
- 3 $\|\hat{Z}\|_* \leq c$ (**Nuclear norm** = sum of singular values)
- 4 $\hat{Z} = UDV^T$ with $\Phi(u_j) \leq c_1, \Phi_2(v_k) \leq c_2$ (Penalized SVD)
- 5 $\hat{Z} = LR^T$ with $\Phi_1(L) \leq c_1, \Phi_2(R) \leq c_2$ (Max-margin matrix factorization)
- 6 $\hat{Z} = L + S$ with $\Phi_1(L) \leq c_1, \Phi_2(S) \leq c_2$ (Additive matrix decomposition)

The Frobenius Norm

- **[Squared] Frobenius norm** = “sum of squares” (for a matrix)
- For a matrix $M \in \mathbf{R}^{m \times n}$, the Frobenius norm of M is

$$\|M\|_F^2 = \sum_{i=1}^m \sum_{j=1}^n m_{ij}^2.$$

- The ℓ_2 -loss at the matrix level is often written as

$$\|Z - \hat{Z}\|_F^2 = \sum_{i=1}^m \sum_{j=1}^n (z_{ij} - \hat{z}_{ij})^2$$

General Framework

- Given $Z \in \mathbf{R}^{m \times n}$,
- Find $\hat{Z} \in \mathbf{R}^{m \times n}$ solving the following optimization problem:

$$\begin{array}{ll} \text{minimize} & \|Z - \hat{Z}\|_F^2 \\ \text{such that} & \Phi(\hat{Z}) \leq c. \end{array}$$

- Modifications required when
 - Z isn't fully observed, or
 - if we want to apply Φ to factors of \hat{Z}

Matrix Notation

- Consider

$$M = \begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \\ m_{31} & m_{33} \end{pmatrix}$$

- Can also write M in terms of its column vectors: $M = (m_{.1}, m_{.2})$.

Matrix-Vector Product

- Consider Mv , where $v = (v_1, v_2)^T$ is a column vector:

$$\begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \\ m_{31} & m_{33} \end{pmatrix} \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} = \begin{pmatrix} v_1 m_{11} + v_2 m_{12} \\ v_1 m_{21} + v_2 m_{22} \\ v_1 m_{31} + v_2 m_{33} \end{pmatrix} = v_1 \begin{pmatrix} m_{11} \\ m_{21} \\ m_{31} \end{pmatrix} + v_2 \begin{pmatrix} m_{12} \\ m_{22} \\ m_{33} \end{pmatrix}$$

- Product is a linear combination of the columns.
- Can write as

$$Mv = v_1 m_{.1} + v_2 m_{.2}$$

Matrix-Matrix Product

- Multiplying 2 matrices is a matrix-vector multiply for each column in product

$$\begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \\ m_{31} & m_{33} \end{pmatrix} \begin{pmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \end{pmatrix} \\ = \left(\begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \\ m_{31} & m_{33} \end{pmatrix} \begin{pmatrix} a_{11} \\ a_{21} \end{pmatrix} ; \begin{pmatrix} m_{11} & m_{12} \\ m_{21} & m_{22} \\ m_{31} & m_{33} \end{pmatrix} \begin{pmatrix} a_{12} \\ a_{22} \end{pmatrix} \right)$$

- From this, easy to see certain results on rank coming up...

Rank of a Matrix

- Let $M \in \mathbf{R}^{m \times n}$ be an $m \times n$ matrix.
- Many equivalent definitions of $\text{rank}(M)$.
- $\text{rank}(M)$ is
 - max number of linearly independent columns
 - max number of linearly independent rows
 - dimension of column space
 - dimension of row space
 - number of non-zero singular values in SVD
- Largest possible rank for M ?

Rank of an outer product

- Suppose we have column vectors $v = (v_1, v_2)^T$ and $w = (w_1, w_2, w_3)^T$.
- Their **outer product** is

$$\begin{aligned}vw^T &= \begin{pmatrix} v_1 \\ v_2 \end{pmatrix} (w_1 \quad w_2 \quad w_3) \\&= \begin{pmatrix} v_1 w_1 & v_1 w_2 & v_1 w_3 \\ v_2 w_1 & v_2 w_2 & v_2 w_3 \end{pmatrix} \\&= (w_1 v \quad w_2 v \quad w_3 v)\end{aligned}$$

- Every column of vw^T is a multiple of v .
- What is the rank of vw^T ?

Every Rank 1 Matrix is an Outer Product

- Suppose $M \in \mathbf{R}^{m \times n}$ is a rank 1 matrix.
- Claim: $\exists w \in \mathbf{R}^m, x \in \mathbf{R}^n$ s.t. $M = wx^T$.
- Proof:
 - Rank 1 implies every column is multiple of first column.
 - Let w =first column of M .
 - Take v to be vector of multiples of w required.

SVD for Rank 1 Matrix

- The factorization

$$M = \sigma_1 uv^T,$$

where $\sigma_1 > 0$ and $\|u\| = \|v\| = 1$ is the **singular value decomposition (SVD)**.

- Easy derivation from $M = wx^T$.

Rank of matrix product

- Suppose $\text{rank}(M) = r$ and $\text{rank}(N) = s$.
- What is largest possible rank of MN ?
- Approach:
 - Every column of MN is a linear combination of columns of M .
 - So $\text{rank}(MN) \leq r$.
- Using $\text{rank}(A) = \text{rank}(A^T)$,
 - can apply same argument to $NM \implies \text{rank}(NM) \leq s$.
- So

$$\text{rank}(MN) \leq \min(\text{rank}(M), \text{rank}(N)).$$

Low Rank Matrices

- $M \in \mathbf{R}^{m \times n}$ is **low rank** if $\text{rank}(M) \ll \min(m, n)$.
- (M has relatively few independent columns.)
- Suppose M is 5×4 and $\text{rank}(M) = 2$, then we can **factorize** as:

$$\underbrace{\begin{pmatrix} m_{11} \\ \vdots \\ m_{54} \end{pmatrix}}_{5 \times 4} = \underbrace{\begin{pmatrix} a_{11} \\ \vdots \\ a_{52} \end{pmatrix}}_{5 \times 2} \underbrace{\begin{pmatrix} b_{11} & \dots & b_{24} \end{pmatrix}}_{2 \times 5}$$

- Proof: Should be clear from discussion above.

Rank-Constrained Matrix Approximation

- Given $Z \in \mathbf{R}^{m \times n}$,
- Find $\hat{Z} \in \mathbf{R}^{m \times n}$ solving the following optimization problem:

$$\begin{array}{ll} \text{minimize} & \|Z - \hat{Z}\|_F^2 \\ \text{such that} & \text{rank}(\hat{Z}) \leq c, \end{array}$$

for $c \in \{1, 2, \dots\}$

- The solution to this problem is immediate from the **singular value decomposition** (SVD)

Singular Value Decomposition (SVD)

- If $M \in \mathbb{R}^{m \times n}$ has rank r , then its **singular value decomposition** is

$$M = \sum_{i=1}^r \sigma_i u_i v_i^T,$$

where

- $\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_r > 0$ are called the **singular values**
- u_1, \dots, u_r are **orthonormal** and are called the **left singular vectors**
- v_1, \dots, v_r are **orthonormal** and are called the **right singular vectors**
- **THEOREM:** Every matrix has a singular value decomposition.

SVD: Rank r Approximation Theorem

- Given $Z \in \mathbf{R}^{m \times n}$, find $\hat{Z} \in \mathbf{R}^{m \times n}$ solving the following:

$$\begin{array}{ll} \text{minimize} & \|Z - \hat{Z}\|_F^2 \\ \text{such that} & \text{rank}(\hat{Z}) \leq c, \end{array}$$

for $c \in \{1, 2, \dots, \text{rank}(Z)\}$.

- THEOREM:** This optimization problem is solved by

$$\hat{Z} = \sum_{i=1}^c \sigma_i u_i v_i^T,$$

where σ 's, u 's, and v 's are given by the SVD for Z :

$$Z = \sum_{i=1}^r \sigma_i u_i v_i^T$$

The Netflix Problem

- Partially observed ratings matrix

Movie Ratings	Zora	Sophie	Jordan	Ernie	Christie
Harold and Kumar Escape..	8	4	?	?	4
Ted	?	?	8	10	4
Straight Outta Compton	8	10	?	?	6

Taken from <https://analyticsweek.com/content/movie-recommendations-how-does-netflix-do-it-a-9-step-coding-intuitive-guide-into-collaborative-filtering/>

The Netflix Problem

- How do we fill in the missing entries of the ratings matrix?
- Let's generalize the problem a bit:
 - m users
 - n movies
 - **ratings matrix** $Z \in \mathbf{R}^{m \times n}$ (an $m \times n$ matrix)
- We observe some entries of Z – how to estimate the rest?

Partial Observations (Notation)

- z_{ij} : rating for user i and movie j
- Observed entries $\mathcal{D} = \{z_{1,2} = 3, z_{2,4} = 0, z_{4,2} = 1, z_{6,2} = 5\}$
 - (This is our **training data**.)
- Ω : set of entries that we observe, e.g.

$$\Omega = \{(1, 2), (2, 4), (4, 2), (6, 2)\}$$

- Mean observed rating:

$$\mu = \frac{1}{|\Omega|} \sum_{(i,j) \in \Omega} z_{ij} = \frac{9}{4}$$

The Prediction Problem

- Given a new user/movie pair (i, j) ,
 - give a prediction \hat{z}_{ij} for the unknown z_{ij} .
- Alternatively,
 - produce a full matrix $\hat{Z} = (\hat{z}_{ij}) \in \mathbf{R}^{m \times n}$ all at once.
- Equivalent, just two different ways to think about it.

Performance Measure


- We predict \hat{z}_{ij} .
- Later we get the actual z_{ij} .
- How to evaluate how well we did?
- Need a **loss function**.
- Netflix used ℓ_2 -loss (i.e. square loss):

$$\ell(z_{ij}, \hat{z}_{ij}) = (z_{ij} - \hat{z}_{ij})^2$$

- We'll focus on square loss.


More Matrix Multiplication

- Sometimes we need the most vanilla description of matrix mult:



$$X_{32} = (\mathbf{a}, \mathbf{b}, \mathbf{c}) \cdot (\mathbf{x}, \mathbf{y}, \mathbf{z}) = \mathbf{a} * \mathbf{x} + \mathbf{b} * \mathbf{y} + \mathbf{c} * \mathbf{z}$$

Netflix Matrix Factorization ("Latent Factor Model")



$$x_{32} = (\mathbf{a}, \mathbf{b}, \mathbf{c}) \cdot (\mathbf{x}, \mathbf{y}, \mathbf{z}) = \mathbf{a} * \mathbf{x} + \mathbf{b} * \mathbf{y} + \mathbf{c} * \mathbf{z}$$

- Factorization tells us $\text{rank}(r) \leq 3$. (clear?)
- x_{32} is rating for user 3 of movie 2.
- Interpretation: we have 3 movie **categories** (called **factors** in more generic contexts)
- q_3 . gives user 3's weightings to each **category**
- $p_{.2}$ gives how much each movie belongs to **category 2**

Alex Lin's "Introduction to Matrix Factorization Methods Collaborative Filtering"

Netflix Matrix Factorization: Objective Function

- Find $\hat{Z} = QP^T$ where Q is $m \times r$ and P is $n \times r$.
- Training loss is

$$\frac{1}{|\Omega|} \sum_{(i,j) \in \Omega} (z_{ij} - q_i \cdot p_j)^2$$

- How many parameters? $r(m+n)$
- Minimize this loss using
 - SGD or
 - Alternating least squares

Put in some regularization

- Training loss is

$$\frac{1}{|\Omega|} \sum_{(i,j) \in \Omega} (z_{ij} - q_{i \cdot p \cdot j})^2 + \sum_{i=1}^m \sum_{j=1}^r (\lambda_1 |q_{ij}| + \lambda_2 q_{ij}^2)$$

- How many parameters? $r(m+n)$
- First pass: minimize this loss using SGD.

Put in some bias terms

- Training loss is

$$\frac{1}{|\Omega|} \sum_{(i,j) \in \Omega} (z_{ij} - [q_i \cdot p_j + u_i + m_j + c])^2 + \sum_{i=1}^m \sum_{j=1}^r (\lambda_1 |q_{ij}| + \lambda_2 q_{ij}^2)$$

- user bias term: u_i
- movie bias term: m_j