

# Recitation 9

## Gradient Boosting

Brett Bernstein

CDS at NYU

March 30, 2017

# Intro Question

## Question

Suppose 10 different meteorologists have produced functions  $f_1, \dots, f_{10} : \mathbb{R}^d \rightarrow \mathbb{R}$  that forecast tomorrow's noon-time temperature using the same  $d$  features. Given a validation set containing 1000 data points  $(x_i, y_i) \in \mathbb{R}^d \times \mathbb{R}$  of similar forecast situations, describe a method to forecast tomorrow's noon-time temperature. Would you use boosting, bagging or neither?

# Intro Solution

## Solution

Let  $\hat{x}_i = (x_i, f_1(x_i), \dots, f_{10}(x_i)) \in \mathbb{R}^{d+10}$ . Then use any fitting method you like to produce an aggregate decision function  $f : \mathbb{R}^{d+10} \rightarrow \mathbb{R}$ . This method is sometimes called stacking.

- 1 This isn't bagging - we didn't generate bootstrap samples and learn a decision function on each of them.
- 2 This isn't boosting - boosting learns decision functions on varying datasets to produce an aggregate classifier.

# Different Ensembles

- 1 Parallel ensemble: each base model is fit independently of the other models. Examples are bagging and stacking.
- 2 Sequential ensemble: each base model is fit in stages depending on the previous fits. Examples are AdaBoost and Gradient Boosting.

# AdaBoost Review

- 1 Recall that a learner, or learning algorithm take a dataset as input and produces a decision function in some hypothesis space.

## Question

Suppose we had a learner that given a dataset, and a weighting (importance) scheme on that dataset, would produce a classifier  $h$  that has lower than .5 loss using the weighted 0 – 1 loss:

$$\frac{1}{n} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq h(x_i)) \leq \gamma < .5.$$

Can we use this learner to create an ensemble that makes accurate predictions?

# AdaBoost Review

- 1 Recall that a learner, or learning algorithm take a dataset as input and produces a decision function in some hypothesis space.

## Question

Suppose we had a learner that given a dataset, and a weighting (importance) scheme on that dataset, would produce a classifier  $h$  that has lower than .5 loss using the weighted 0 – 1 loss:

$$\frac{1}{n} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq h(x_i)) \leq \gamma < .5.$$

Can we use this learner to create an ensemble that makes accurate predictions?

- We saw that AdaBoost solves this problem.
- Can get around weighted loss functions using sampling trick.

# Additive Models

- ① Additive models over a base hypothesis space  $\mathbb{H}$  take the form

$$\mathcal{F} = \left\{ f(x) = \sum_{m=1}^M \nu_m h_m(x) \mid h_m \in \mathbb{H}, \nu_m \in \mathbb{R} \right\}.$$

- ② Since we are taking linear combinations, we assume the  $h_m$  functions take values in  $\mathbb{R}$  or some other vector space.
- ③ Empirical risk minimization over  $\mathcal{F}$  tries to find

$$\arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- ④ This in general is a difficult task, as the number of base hypotheses  $M$  is unknown, and each base hypothesis  $h_m$  ranges over all of  $\mathbb{H}$ .

# Forward Stagewise Additive Modeling (FSAM)

The FSAM method fits additive models using the following (greedy) algorithmic structure:

- ① Initialize  $f_0 \equiv 0$ .
- ② For stage  $m = 1, \dots, M$ :
  - ① Choose  $h_m \in \mathbb{H}$  and  $\nu_m \in \mathbb{R}$  so that

$$f_m = f_{m-1} + \nu_m h_m$$

has the minimum empirical risk.

- ② The function  $f_m$  has the form

$$f_m = \nu_1 h_1 + \dots + \nu_m h_m.$$



# Forward Stagewise Additive Modeling (FSAM)

The FSAM method fits additive models using the following (greedy) algorithmic structure:

- 1 Initialize  $f_0 \equiv 0$ .
- 2 For stage  $m = 1, \dots, M$ :
  - 1 Choose  $h_m \in \mathbb{H}$  and  $\nu_m \in \mathbb{R}$  so that

$$f_m = f_{m-1} + \nu_m h_m$$

has the minimum empirical risk.

- 2 The function  $f_m$  has the form

$$f_m = \nu_1 h_1 + \dots + \nu_m h_m.$$

- When choosing  $h_m, \nu_m$  during stage  $m$ , we must solve the minimization

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathbb{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

# Gradient Boosting

- 1 Can we simplify the following minimization problem:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathbb{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

- 2 What if we linearize the problem and take a step along the steepest descent direction?
- 3 Good idea, but how do we handle the constraint that  $h$  is a function that lies in  $\mathbb{H}$ , the base hypothesis space?

# Gradient Boosting

- ① Can we simplify the following minimization problem:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathbb{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

- ② What if we linearize the problem and take a step along the steepest descent direction?
- ③ Good idea, but how do we handle the constraint that  $h$  is a function that lies in  $\mathbb{H}$ , the base hypothesis space?
- ④ First idea: since we are doing empirical risk minimization, we only care about the values  $h$  takes on the training set. Thus we can think of  $h$  as a vector  $(h(x_1), \dots, h(x_n))$ .
- ⑤ Second idea: first compute unconstrained steepest descent direction, and then constrain (project) onto possible choices from  $\mathbb{H}$ .

# Gradient Boosting Machine

- ① Initialize  $f_0 \equiv 0$ .
- ② For stage  $m = 1, \dots, M$ :
  - ① Compute the steepest descent direction (also called *pseudoresiduals*):

$$r_m = -(\partial_2 \ell(y_1, f_{m-1}(x_1)), \dots, \partial_2 \ell(y_n, f_{m-1}(x_n))).$$

- ② Find the closest base hypothesis (using Euclidean distance):

$$h_m = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

- ③ Choose fixed step size  $\nu_m \in (0, 1]$  or line search:

$$\nu_m = \arg \min_{\nu \geq 0} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h_m(x_i)).$$

- ④ Take the step:

$$f_m(x) = f_{m-1}(x) + \nu_m h_m(x).$$

# Gradient Boosting Machine

- 1 Each stage we need to solve the following step:

$$h_m = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

How do we do this?

# Gradient Boosting Machine

- 1 Each stage we need to solve the following step:

$$h_m = \arg \min_{h \in \mathbb{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

How do we do this?

- 2 This is a standard least squares regression task on the “mock” dataset

$$\mathcal{D}^{(m)} = \{(x_1, (r_m)_1), \dots, (x_n, (r_m)_n)\}.$$

- 3 We assume that we have a learner that (approximately) solves least squares regression over  $\mathbb{H}$ .

# Gradient Boosting Comments

- 1 The algorithm above is sometimes called AnyBoost or Functional Gradient Descent.
- 2 The most commonly used base hypothesis space is small regression trees (HTF recommends between 4 and 8 leaves).

# Practice With Different Loss Functions

## Question


Explain how to perform gradient boosting with the following loss functions:

- 1 Square loss:  $\ell(y, a) = (y - a)^2/2$ .
- 2 Absolute loss:  $\ell(y, a) = |y - a|$ .
- 3 Exponential margin loss:  $\ell(y, a) = e^{-ya}$ .



# Demonstration Using Decision Stumps

Below is an example of a decision stump for functions  $h : \mathbb{R} \rightarrow \mathbb{R}$ .



# Demonstration Using Decision Stumps

Below is the dataset we will use.

