

Recitation 9

Gradient Boosting

Brett Bernstein

CDS at NYU

March 29, 2017

Intro Question

Question

Suppose 10 different meteorologists have produced functions $f_1, \dots, f_{10} : \mathbb{R}^d \rightarrow \mathbb{R}$ that forecast tomorrow's noon-time temperature using the same d features. Given a validation set containing 1000 data points $(x_i, y_i) \in \mathbb{R}^d \times \mathbb{R}$ of similar forecast situations, describe a method to forecast tomorrow's noon-time temperature. Would you use boosting, bagging or neither?

Intro Solution

Solution

Let $\hat{x}_i = (x_i, f_1(x_i), \dots, f_{10}(x_i)) \in \mathbb{R}^{d+10}$. Then use any fitting method you like to produce an aggregate decision function $f : \mathbb{R}^{d+10} \rightarrow \mathbb{R}$. This method is sometimes called stacking.

- 1 This isn't bagging - we didn't generate bootstrap samples and learn a decision function on each of them.
- 2 This isn't boosting - boosting learns decision functions on varying datasets to produce an aggregate classifier.

Different Ensembles

- 1 Parallel ensemble: each base model is fit independently of the other models. Examples are bagging and stacking.
- 2 Sequential ensemble: each base model is fit in stages depending on the previous fits. Examples are AdaBoost and Gradient Boosting.

AdaBoost Review

- 1 Recall that a learner, or learning algorithm take a dataset as input and produces a decision function in some hypothesis space.

Question

Suppose we had a learner that given a dataset, and a weighting (importance) scheme on that dataset, would produce a classifier h that has lower than .5 loss using the weighted 0 – 1 loss:

$$\frac{1}{n} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq h(x_i)) \leq .5.$$

Can we use this learner to create an ensemble that makes accurate predictions?

AdaBoost Review

- 1 Recall that a learner, or learning algorithm take a dataset as input and produces a decision function in some hypothesis space.

Question

Suppose we had a learner that given a dataset, and a weighting (importance) scheme on that dataset, would produce a classifier h that has lower than .5 loss using the weighted 0 – 1 loss:

$$\frac{1}{n} \sum_{i=1}^n w_i \mathbf{1}(y_i \neq h(x_i)) \leq .5.$$

Can we use this learner to create an ensemble that makes accurate predictions?

- We saw that AdaBoost solves this problem.
- Can get around weighted loss functions using sampling trick.

Additive Models

- ① Additive models over a base hypothesis space \mathcal{H} take the form

$$\mathcal{F} = \left\{ f(x) = \sum_{m=1}^M \nu_m h_m(x) \mid h_m \in \mathcal{H}, \nu_m \in \mathbb{R} \right\}.$$

- ② Since we are taking linear combinations, we assume the h_m functions take values in \mathbb{R} or some other vector space.
- ③ Empirical risk minimization over \mathcal{F} tries to find

$$\arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n \ell(y_i, f(x_i)).$$

- ④ This in general is a difficult task, as the number of base hypotheses M is unknown, and each base hypothesis h_m ranges over all of \mathcal{H} .

Forward Stagewise Additive Modeling (FSAM)

The FSAM method fits additive models using the following (greedy) algorithmic structure:

- ① Initialize $f_0 \equiv 0$.
- ② For stage $m = 1, \dots, M$:
 - ① Choose $h_m \in \mathcal{H}$ and $\nu_m \in \mathbb{R}$ so that

$$f_m = f_{m-1} + \nu_m h_m$$

has the minimum empirical risk.

- ② The function f_m has the form

$$f_m = \nu_1 h_1 + \dots + \nu_m h_m.$$

Forward Stagewise Additive Modeling (FSAM)

The FSAM method fits additive models using the following (greedy) algorithmic structure:

- 1 Initialize $f_0 \equiv 0$.
- 2 For stage $m = 1, \dots, M$:
 - 1 Choose $h_m \in \mathcal{H}$ and $\nu_m \in \mathbb{R}$ so that

$$f_m = f_{m-1} + \nu_m h_m$$

has the minimum empirical risk.

- 2 The function f_m has the form

$$f_m = \nu_1 h_1 + \dots + \nu_m h_m.$$

- When choosing h_m, ν_m during stage m , we must solve the minimization

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathcal{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

Gradient Boosting

- ① Can we simplify the following minimization problem:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathcal{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

- ② What if we linearize the problem and take a step along the steepest descent direction?
- ③ Good idea, but how do we handle the constraint that h is a function that lies in \mathcal{H} , the base hypothesis space?

Gradient Boosting

- ① Can we simplify the following minimization problem:

$$(\nu_m, h_m) = \arg \min_{\nu \in \mathbb{R}, h \in \mathcal{H}} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h(x_i)).$$

- ② What if we linearize the problem and take a step along the steepest descent direction?
- ③ Good idea, but how do we handle the constraint that h is a function that lies in \mathcal{H} , the base hypothesis space?
- ④ First idea: since we are doing empirical risk minimization, we only care about the values h takes on the training set. Thus we can think of h as a vector $(h(x_1), \dots, h(x_n))$.
- ⑤ Second idea: first compute unconstrained steepest descent direction, and then constrain (project) onto possible choices from \mathcal{H} .

Gradient Boosting Machine

- ① Initialize $f_0 \equiv 0$.
- ② For stage $m = 1, \dots, M$:
 - ① Compute the steepest descent direction (also called *pseudoresiduals*):

$$r_m = -(\partial_2 \ell(y_1, f_{m-1}(x_1)), \dots, \partial_2 \ell(y_n, f_{m-1}(x_n))).$$

- ② Find the closest base hypothesis (using Euclidean distance):

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

- ③ Choose fixed step size $\nu_m \in (0, 1]$ or line search:

$$\nu_m = \arg \min_{\nu \geq 0} \sum_{i=1}^n \ell(y_i, f_{m-1}(x_i) + \nu h_m(x_i)).$$

- ④ Take the step:

$$f_m(x) = f_{m-1}(x) + \nu_m h_m(x).$$

Gradient Boosting Machine

- 1 Each stage we need to solve the following step:

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

How do we do this?

Gradient Boosting Machine

- 1 Each stage we need to solve the following step:

$$h_m = \arg \min_{h \in \mathcal{H}} \sum_{i=1}^n ((r_m)_i - h(x_i))^2.$$

How do we do this?

- 2 This is a standard least squares regression task on the “mock” dataset

$$\mathcal{D}^{(m)} = \{(x_1, (r_m)_1), \dots, (x_n, (r_m)_n)\}.$$

- 3 We assume that we have a learner that (approximately) solves least squares regression over \mathcal{H} .

Gradient Boosting Comments

- 1 The algorithm above is sometimes called AnyBoost or Functional Gradient Descent.
- 2 The most commonly used base hypothesis space is small regression trees (HTF recommends between 4 and 8 leaves).

Practice With Different Loss Functions


Question

Explain how to perform gradient boosting with the following loss functions:

- 1 Square loss: $\ell(y, a) = (y - a)^2/2$.
- 2 Absolute loss: $\ell(y, a) = |y - a|$.
- 3 Exponential margin loss: $\ell(y, a) = e^{-ya}$.

Demonstration Using Decision Stumps

Below is an example of a decision stump for functions $h : \mathbb{R} \rightarrow \mathbb{R}$.



Demonstration Using Decision Stumps

Below is the dataset we will use.

