

```
In [1]: from IPython import display
def hide_code_in_slideshow():
    import os
    uid = os.urandom(8).encode("hex")
    html = """<div id="%s"></div>
<script type="text/javascript">
    $(function(){
        var p = $("#%s");
        if (p.length==0) return;

        while (!p.hasClass("cell")) {
            p=p.parent();

            if (p.prop("tagName") == "body") return;
        }
        var cell = p;
        cell.find(".input").addClass("hide-in-slideshow")
    });
</script>""" % (uid, uid)
    display.display_html(html, raw=True)
```

# Feature Importance outline

## Feature importance

- Introduction and motivation for discussing feature importance:
  - Define feature importance (using a fuzzy definition -- rank features by how much they "*contribute*" to the prediction function).
  - Feature importance measures are frequently reported, in part because many packages make them easy to compute/visualize.
    - Show google image search results
  - They can be potentially misleading. Important to understand potential pitfalls.
  - Utility of feature importance:
    - Feature selection -- use votes example from original Breiman Random Forests paper.
    - Sanity checking models/detecting leakage (TODO: see if I can find an example where feature importance show some feature unexpectedly popping since due to leakage).
- Framing feature importance methods:
  - Divide into two classes:
    - Methods that exploit the structure of the prediction function or learning algorithm (call these function-specific).
    - Methods that are agnostic (i.e. can be applied to any black box prediction function).
- Function-specific methods:
  - Mean decrease impurity:
    - Frame running example:
      - Introduce the iris dataset and classification problem.
      - graphviz some shallow tree on the iris dataset (see below).
      - **Ask students** which features seem "important" as a turn and talk. Have some subset share answers without explanation.
    - Introduce MDE:

$$Imp(X_m) = \frac{1}{N_T} \sum_T \sum_{t \in T: v(s_t) = X_m} p(t) \Delta i(s_t, t)$$

- Show from scratch implimentation (see lab\_outline\_and\_mean\_decrease\_impurity.ipynb).
    - Show correspondance with sklearn's implementation.
    - **Ask students** what other classifiers we could extend this too (Random Forests/GBMs) and how (averaging).
    - Show RF importance for the iris problem.
    - Say there are variants:
      - Example: CART implementation exposes a knob to add weight for use of a feature as a surrogate split.
  - Absolute coefficients in a linear model:
    - **Ask students** how we could determine feature importance in a linear model.
    - Show some example (TODO) of extracting absolute  $\beta$ s from a model.
    - **Ask students:** How would this interact with preprocessing?
- Model agnostic methods:
  - Frame the general problem for global feature importance:

- Another framing of feature importance -- we are trying to identify how much each feature independently contributes to risk reduction in the learned model.
- **Ask students:** How could we do this for an arbitrary black box prediction function?
- **Hint (fragment slide):** Imagine you were tasked with constructing a synthetic dataset with an *unimportant* feature -- how could you construct it? Given this insight, what operation could be applied to a feature that would be expected to have (a) no impact on empirical risk for unimportant features, and (b) decrease performance for important features?
- Permutation importance:
  - Give Breimans original OOB permutation feature importance.
  - Say that the OOB version exploits bootstrapping in RFs, but can just use a validation set.
  - Show example (using ELI5 implementation) over the iris problem.
  - **Ask students:** This method permute features to get a measure of feature importance. Let's extend this concept a bit -- how could permuting the target be useful in interrogating the performance of a model?
  - <http://jmlr.csail.mit.edu/papers/volume11/ojala10a/ojala10a.pdf> (<http://jmlr.csail.mit.edu/papers/volume11/ojala10a/ojala10a.pdf>).
  - For fun show this using `permutation_test_score` for the iris data.
- Pitfalls:
  - Correlated features
    - Show the GBM feature importance for duplicated or highly correlated features (TODO: think about which is better for this purpose) (from `lab_outline_and_mean_decrease_impurity.ipynb`).
    - **Ask students:** Interpret this plot. What can we say about the features/model?
    - Show same plot for enet, and ask same question -- hopefully students will recall that (at least in the context of linear models) our conversations about feature correlation.
    - Show the elastic net bound result as review, and primer to think about correlation issues.
    - Show correlation matrix/matrices (duplicated and/or highly correlated), and known data generating process.
    - **Ask students:** Given what we know about (i) linear methods, and (ii) trees, what does this show us about potential pitfalls in interpreting feature importance plots?
    - **Ask students:** How would this impact SelectFromModel feature selection procedures (ala sklearn) which select the most important features from a model based on either absolute coefficients or feature importance scores? Why might we want to use RFE?

## PDPs

- Motivation and introduction to partial dependence plots
  - Imagine we have a subset of features we think are important.
    - This could be based on "feature importance" scores or coefficients, or other reasons (i.e. prior knowledge/research questions/etc.).
  - We may want to dig deeper to explain the relationship between our predictions and these features.
    - **Ask students:** Why? If we have MDI feature importance scores in hand, what do we not know?
      - Directionality:
        - Note in reality we probably have complex, non-monotonic relationships
        - However even if we have a monotonic relationship, or even linear, our MDI wouldn't give any indication on the direction -- obviously weights in a linear model would.

- Partial dependence plots let us dig deeper and visualize the dependence of our prediction function (i.e. predict/predict\_proba) on one or two of these features. Can't easily go higher-dimensional (Why not? Can't visualize).
- Show basic math (marginal and empirical estimate -- see <https://arxiv.org/pdf/1309.6392.pdf> (<https://arxiv.org/pdf/1309.6392.pdf>) for useful framing).
- PDP examples:
  - Show sklearn example (see partial\_dependence.ipynb).
  - **Ask students:** What does this plot show? How do we interpret it?
- PDP pitfalls:
  - Show the flat partial dependence plot from (<https://arxiv.org/pdf/1309.6392.pdf> (<https://arxiv.org/pdf/1309.6392.pdf>)) reproduced in partial\_dependence.ipynb.
  - **Ask students:** What does this plot show?
  - Follow up with big reveal:
    - Show data generating process and X plot.
  - **Ask students:** What does this teach us? When might PDPs fail?
  - Share "not too strong" quote from Friedman's original work (see partial\_dependence.ipynb).

## Additional material

- If 100 minutes allows (what is your opinion, David?), the next steps would be:
  - Generate the ICE plot for the X data (figure 2 in <https://arxiv.org/pdf/1309.6392.pdf> (<https://arxiv.org/pdf/1309.6392.pdf>))
  - Take this as a launching point for point out that our progression has been from:
    - Global measures (feature importance) that don't necessarily describe structure (even directionality) of a relationship between a feature and target.
    - Marginalized measures showing average relationship between one or two features and the target.
    - Data-instance specific feature importance plots showing the importance of a feature, for each data instance, fixing all other feature values.
  - We can follow this thread it's conclusion -- attempting to develop model explanations that show the relative contribution of each feature to predictions locally (i.e. in the neighborhood of a specific instance):
    - LIME paper (<https://arxiv.org/abs/1602.04938> (<https://arxiv.org/abs/1602.04938>)) as one example.

# Feature Importance

Ben Jakubowski  
 NYU  
 May 1st, 2018

## Learning objectives

### ***Feature Importance***

- (Informally) define feature importance, and explain why exploring and presenting feature importance can be both (i) useful, and (ii) potentially misleading.
- Define and describe several feature importance methods that exploit the structure of the learning algorithm or learned prediction function.
- Describe a prediction-function-agnostic method for generating feature importance scores.
- Describe the limitations of these feature importance measures and understand cases where they "fail".

### ***Partial Dependence Plots***

- Describe how partial dependence plots are generated.
- Describe the limitations of these plots, and some potential remedies.

## Defining feature importance

- Problem setup -- we have some data  $(x, y) \in (\mathbb{R}^d, \mathcal{Y})$ , and a learned prediction function  $f$ .
- A feature importance method can be loosely understood as a function that maps each feature onto some score.
- These scores rank features by how much they "*contribute*" to the prediction function  $f$ .
  - Here "*contribute*" is defined separately for each method.
  - In general, *feature importance* is not consistently or rigorously defined.

## Why we need to discuss feature importance

- It is easy to compute feature importance from many ML libraries.



feature importance



All **Images** News Videos Shopping More

Settings Tools

View saved SafeSearch

forward

sequential

data mining

machine learning

importance

filter

variable














attribute

subset

classification

text classification

feature extraction



## Why we need to discuss feature importance


- It is easy to compute feature importance from many ML libraries.
- This means feature importance bar charts show up all over, and it is important to understand:
  - How feature importance can be computed.
  - Potential pitfalls in over interpreting these scores.

## Utility of feature importance scores

- **Question:** Why might feature importance scores be useful?
- Use for sanity checking a model -- do the features that are important seem reasonable, or do they suggest something strange is going on under the hood (i.e. leakage)?
- Share with clients/bosses to build trust and get buy-in on the prediction function.
- Use for feature selection -- can extract the top  $K$  features and retrain over this subset (ex: `sklearn.feature_selection.SelectFromModel`)

## Feature selection example

- Breiman's original Random Forests paper showed feature importance for classification of political party based on 16 votes<sup>[1]</sup>  
(<https://link.springer.com/content/pdf/10.1023/A:1010933404324.pdf>).



*Figure 6.* Measure of variable importance—votes data.

- 'We reran this data set using only variable 4. The test set error is 4.3%, about the same as if all variables were used.'



## First example/introducing the issues

- Consider the following feature importance plot for a regression problem with

$$\mathcal{Y} = \mathcal{A} = \mathbb{R}, X = \mathbb{R}^{11}$$


```
In [2]: hide_code_in_slideshow()
from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
from sklearn.linear_model import ElasticNetCV
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
sns.set_context('notebook', font_scale=1.5)
%matplotlib inline

X,y = make_regression(n_samples=1000, n_features=10,
                     n_informative=3,random_state=1234,
                     shuffle=False, noise=10)
X = np.append(X, X[:,0], axis=1)
train_X, test_X, train_y, test_y = train_test_split(X,y,
                                                    test_size=0.2,
                                                    random_state=1235)

enet = ElasticNetCV(l1_ratio=np.array([.1, .5, .7, .9, .95, .99, 1]),
                   n_alphas=500,
                   normalize=True,
                   selection='random',
                   random_state=1351)
enet.fit(train_X, train_y)

coefs = pd.DataFrame({'Feature': np.arange(len(enet.coef_)),
                     'Importance': np.abs(enet.coef_)})
coefs = coefs.set_index('Feature')\
            .sort_values('Importance', ascending=False)

fig, ax=plt.subplots(figsize=(12,6))
plot = coefs.plot(kind='bar', title='Example feature importance chart', ax=ax)
```



- Question: What can we say about this prediction function, its features, and target?

### First example/introducing the issues (continued)

- Recall the full feature set includes features  $\{0, \dots, 10\}$ .
- Let's compare both (a) test set performance ( $R^2$ ), and (b) feature importance plots, for models trained on:
  - All features
  - Features  $\{0, \dots, 9\}$
  - Features  $\{1, \dots, 9\}$
  - Features  $\{1, \dots, 10\}$

### First example/introducing the issues (continued)

- Recall the full feature set includes features  $\{0, \dots, 10\}$ ...

```

In [3]: from matplotlib import rc
rc('text', usetex=True)

from sklearn.metrics import mean_squared_error
from math import floor

hide_code_in_slideshow()
features_to_test = {
    'All': np.arange(0,11),
    '0-9': np.arange(0,10),
    '1-9': np.arange(1,10),
    '1-10': np.arange(1,11)
}

fig, axarr = plt.subplots(ncols=2, nrows=2, figsize=(12,8))

i = 0
for model_name, features in features_to_test.items():
    col = i % 2
    row = int(floor(i/2))

    _train_X = train_X[:, features]
    _test_X = test_X[:, features]

    enet = ElasticNetCV(l1_ratio=np.array([.1, .5, .7, .9, .95, .99, 1]),
                        n_alphas=1000,
                        normalize=True,
                        selection='random',
                        random_state=1351)
    enet.fit(_train_X, train_y)

    coefs = pd.DataFrame({'Feature': features,
                        'Importance': np.abs(enet.coef_)})
    coefs = coefs.set_index('Feature')\
        .sort_values('Importance', ascending=False)



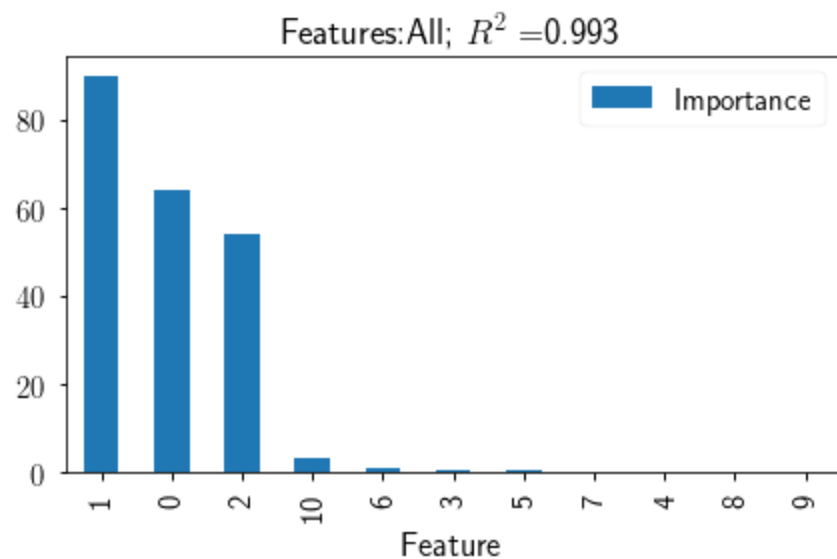

    plot = coefs.plot(kind='bar', title='Typical feature importance chart', ax=axarr[(row,col)])

    title = 'Features:' + model_name + r'; $R^2$ =' + \
        str(round(enet.score(_test_X, test_y), 3))

    axarr[(row,col)].set_title(title)

    i += 1
plt.tight_layout()

```



- What can we learn from this example?

## First example/introducing the issues (continued)

- Here was the setup:

```
X,y = make_regression(n_samples=1000, n_features=10,  
                      n_informative=3, random_state=1234,  
                      shuffle=False, noise=10)  
X = np.append(X, X[:,[0]], axis=1)
```

- Some things to keep in mind:
  - Interactions and correlated/dependent features makes interpretation tricky.
  - Feature importance  $\neq$  dependence of target on feature
  - Feature importance  $\approx$  contribution of feature to prediction function  $f$
  - **This doesn't imply there do not exist other functions  $f'$  where features that are 'unimportant' in  $f$  are 'important'**

## Where we're headed:

- With this in mind, we're going to look at a number of methods for measuring feature importance.
- We can divide these methods into two basic classes:
  - Algorithmic/prediction-function-specific methods: These exploit the structure of the learning algorithm/prediction function to construct some measure of the relative contribution of each feature.
  - Model agnostic methods: These don't make assumptions on the algorithm or structure of the prediction function; can be applied to any black-box predictor.

## Prediction-function-specific: Trees

- Our first algorithm/prediction function specific methods will address decision trees.
- We'll use the Iris Dataset, which is a canonical classification dataset (first used by R.A. Fisher) with:
  - $\mathcal{Y} = \{\text{Setosa, Versicolour, Virginica}\}$  (three varieties of irises).
  - $X = [\text{sepal length, sepal width, petal length, petal width}] \in \mathbb{R}^4$

## Prediction-function-specific: Trees (Continued)

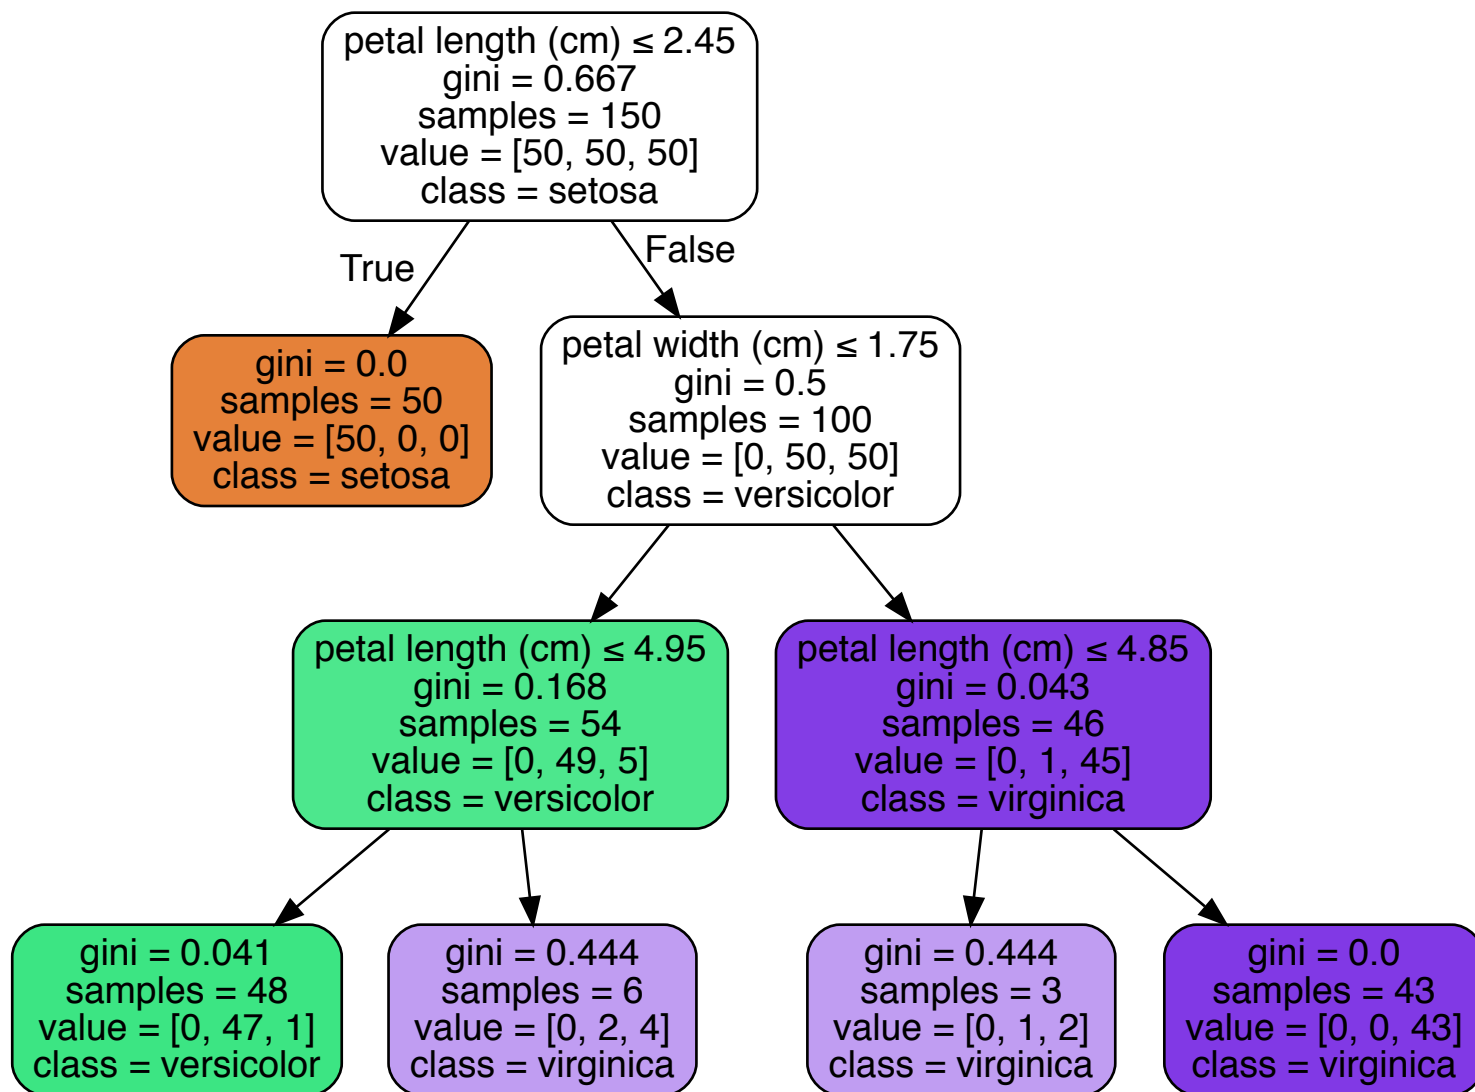
- Let's consider a shallow tree built on this data.

```
In [4]: hide_code_in_slideshow()
from sklearn.datasets import load_iris
from sklearn import tree
import graphviz

iris = load_iris()
clf = tree.DecisionTreeClassifier(max_depth=3)
clf = clf.fit(iris.data, iris.target)

dot_data = tree.export_graphviz(clf, out_file=None,
                                feature_names=iris.feature_names,
                                class_names=iris.target_names,
                                filled=True, rounded=True,
                                special_characters=True)
graph = graphviz.Source(dot_data)
display.display(graph)
```





- **Question:** Which of the feature ([sepal length, sepal width, petal length, petal width]) are most important? Why?

## Prediction-function-specific: Trees (Continued)

### Tree buiding reminders

- Consider node  $t$  in a decision tree built on  $N$  training data instances.
- Let node  $t$  have  $N_t$  node samples.
- We find the split  $s_t$  at node  $t$  such that  $t_L$  and  $t_R$  maximizes the decrease

$$\Delta i(s, t) = i(t) - p_L i(t_L) - p_R i(t_R)$$

where  $p_L$  and  $p_R$  are the probabilities an instance splits left and right, respectively, and  $i(t)$  is some some impurity measure.

## Prediction-function-specific: Trees (Continued)

### Mean decrease impurity

- The **mean decrease impurity**  $imp(X_m)$  for feature  $X_m$  is:

$$imp(X_m) = \sum_{v(s_t)=X_m} p(t) \Delta i(s_t, t)$$

Note  $p(t) = N_t/N$  is the proportion of samples reaching node  $t$  and  $v(s_t)$  is the variable used in split  $s_t$ .

## Prediction-function-specific: Trees (Continued)

$$imp(X_m) = \sum_{v(s_t)=X_m} p(t) \Delta i(s_t, t)$$

- **Question:** When will a feature be considered more important under this metric?
- **Question:** Why is this an algorithm/prediction-function specific method?

## Prediction-function-specific: Trees (Continued)

- Implementation:

```

In [8]: from collections import defaultdict
import numpy as np

def feature_importance_single_tree(tree, feature_names=iris.feature_names):
    # Returns normed feature importance for a single sklearn.tree
    # Note sklearn trees can handle instance weights, so we need
    # to use weighted_n_node_samples

    total_samples = np.sum(tree.weighted_n_node_samples)
    feature_importance = defaultdict(float)

    # Identify leaves as described in sklearn's plot_unveil_tree_structure.html
    is_leaf = (tree.children_right == tree.children_left)
    for ix in range(len(is_leaf)):
        if not is_leaf[ix]:
            impurity = tree.impurity[ix]
            split_feature = tree.feature[ix]
            num_at_node = tree.weighted_n_node_samples[ix]

            # Get left child contribution
            left_child = tree.children_left[ix]
            left_decrease = tree.weighted_n_node_samples[left_child]/num_at_node * \
                            tree.impurity[left_child]

            # Get right child contribution
            right_child = tree.children_right[ix]
            right_decrease = tree.weighted_n_node_samples[right_child]/num_at_node * \
                             tree.impurity[right_child]

            delta = impurity - left_decrease - right_decrease

            feature_importance[feature_names[split_feature]] \
                += num_at_node / total_samples * delta
    norm = np.sum(feature_importance.values())
    feature_importance = {key: val/norm for key, val in feature_importance.items()}
    return feature_importance

```

## Prediction-function-specific: Trees (Continued)

- Implementation vs. sklearn

```
In [9]: imp = feature_importance_single_tree(clf.tree_)
display.display(pd.Series(imp).rename('Importance').to_frame()\
                  .reset_index().rename(columns={'index': 'Feature'}))
```

	Feature	Importance
0	petal length (cm)	0.585616
1	petal width (cm)	0.414384

```
In [10]: display.display(pd.DataFrame({'Feature': iris.feature_names,
                                       'Importance': clf.feature_importances_})\
                        .sort_values('Importance', ascending=False))
```

	Feature	Importance
2	petal length (cm)	0.585616
3	petal width (cm)	0.414384
0	sepal length (cm)	0.000000
1	sepal width (cm)	0.000000

## Prediction-function-specific: Trees (Continued)

- **Question:** How do we extend this to ensembles of trees?

- **Answer:** (Weighted) averages.
  - Ex: Random Forests with  $N_T$  trees:

$$Imp(X_m) = \frac{1}{N_T} \sum_T \sum_{t \in T: v(s_t) = X_m} p(t) \Delta i(s_t, t)$$

## Prediction-function-specific: Trees (Continued)

- Implementation vs. sklearn

```
In [11]: from sklearn.ensemble import RandomForestClassifier
forest = RandomForestClassifier(n_estimators=100) # Not tuned - just example
forest.fit(iris.data, iris.target)
display.display(pd.DataFrame({'Feature': iris.feature_names,
                              'Importance':forest.feature_importances_})\
                .sort_values('Importance', ascending=False))
```

	Feature	Importance
3	petal width (cm)	0.455781
2	petal length (cm)	0.415403
0	sepal length (cm)	0.102569
1	sepal width (cm)	0.026247

```
In [12]: forest_importance = defaultdict(float)
for tree in forest.estimators_:
    tree_importance = feature_importance_single_tree(tree.tree_)

    for key, val in tree_importance.items():
        forest_importance[key] += val
forest_importance = {key:val/len(forest.estimators_)\
                    for key, val in forest_importance.items()}

display.display(pd.Series(forest_importance).rename('Importance').to_frame()\
                .reset_index().rename(columns={'index':'Feature'})\
                .sort_values('Importance', ascending=False))
```

	Feature	Importance
1	petal width (cm)	0.455781
0	petal length (cm)	0.415403
2	sepal length (cm)	0.102569
3	sepal width (cm)	0.026247

## Prediction-function-specific: Trees (Continued)

- Let's think about potential variants/extensions to mean decrease impurity.
- How could we handle surrogate splits?
- Commercial implementation of CART (<https://www.salford-systems.com/support/faq/cart/what-is-variable-importance>) adds contributions of surrogate splits:
  - $\Delta i(s_t, t)$  for primary split variable
  - $p^n \Delta i(s_t, t)$  for the  $n$ 'th surrogate split variable (where  $p \in [0, 1]$  is a user-selected hyperparameter).
- **Question:** Can we think of any other variants or extensions?
- Consider collapsing the importance of dummy-encoded categoricals.

## Prediction-function-specific: Linear models

- **Question:** How can we measure feature importance in a linear model?

$$\text{imp}(X_m) = |w_m|$$

- **Question:** Per usual, how does this notion of importance interact with preprocessing?

## Model-agnostic feature importance:

- We've discussed two methods for algorithm or prediction-function specific feature importance measures:
  - Trees (and their ensembles): Mean Decrease Impurity
  - Linear methods: Absolute weights
- **Question:** How could we determine the importance of input features for an arbitrary black box prediction function?
- Hint 1: Think about defining importance as by measuring how much each feature independently contributes to the performance of the learned model.

- Hint 2: Imagine you were tasked with constructing a synthetic dataset with an unimportant feature -- how could you construct it? Given this insight, what operation could be applied to a feature that would be expected to have (a) no impact on the test set score for unimportant features, and (b) decrease performance for important features?

## Model-agnostic feature importance:

### *Permutation Importance*

- Described by Breiman in the original Random Forests paper (using OOB sample).
- We describe it using an arbitrary held-out test set:
  1. Learn prediction function  $f$  on training data.
  2. Measure the performance of  $f$  on the test set - call this  $s_0$ .
  3. For each feature  $i$  in  $[1, \dots, D]$ :
    - A. Permute just this feature in the test set.
    - B. Pass this permuted test set through  $f$  to obtain a new score  $s_i$ .
  4. Use either  $s_0 - s_i$  (or some appropriate variant on this idea) as feature importance.

## Model-agnostic feature importance:

### *Permutation Importance*

- **Question:** Why bother permuting the feature? Why not just drop it entirely?
- **Question:** Why bother permuting the feature? Why not replace it with some noise, like  $\mathcal{N}(0, 1)$ ?

## Model-agnostic feature importance:

### *Permutation Importance*

- **Question:** Is this any different than the following procedure? If so, how?
  1. Learn prediction function  $f$  on training data.
  2. Measure the performance of  $f$  on the test set - call this  $s_0$ .
  3. For each feature  $i$  in  $[1, \dots, D]$ :
    - A. Drop this feature from both the training and the test set.
    - B. Learn a new prediction function without this feature -- call this function  $f_{\setminus i}$
    - C. Measure the performance of  $f_{\setminus i}$  on the test set to obtain a new score  $s_i$ .
  4. Use either  $s_0 - s_i$  as feature importance.
- **Key difference:** Permutation importance (plus  $|w|$  and MDI) describe the importance of each feature to a particular prediction function  $f$ . We don't claim that a feature with low permutation importance in  $f$  couldn't have high importance in some other  $f'$ .

## Feature importance pitfalls:

- To conclude our discussion of feature importance, we're going to look at an example to understand potential pitfalls.
- We will compare:
  - Permutation importance
  - MDI for a gradient boosting regressor
  - $|w|$  for an elastic net

## Feature importance pitfalls: Problem set up



```
In [14]: import numpy as np
size = 1000
x1 = np.random.uniform(-10,10,size)
z1 = np.random.uniform(-10,10,size)
x2 = z1 + np.random.normal(0,2,size)
x3 = z1 + np.random.normal(0,2,size)
x4 = z1 + np.random.normal(0,2,size)

def actual_function(x1,z1):
    return -1./2.*x1 + z1 + np.random.normal(0,1,len(x1))

y = actual_function(x1,z1)
X = np.stack([x1,x2,x3,x4]).T
```

- **Question:** Which feature is most important (i.e. contributes most to the true conditional mean function)?
- **Question:** Why might we expect this problem to challenge our feature importance methods?

```
In [43]: hide_code_in_slideshow()
corr = pd.DataFrame(X).corr()
corr.index = ['$x_{}$'.format(x) for x in range(4)]
corr.columns = ['$x_{}$'.format(x) for x in range(4)]
display.display(display.Markdown('## Correlation matrix:'))
display.display(corr)
```

## Correlation matrix:

	$x_0$	$x_1$	$x_2$	$x_3$
$x_0$	1.000000	-0.057013	-0.037270	-0.042008
$x_1$	-0.057013	1.000000	0.901286	0.898897
$x_2$	-0.037270	0.901286	1.000000	0.903997
$x_3$	-0.042008	0.898897	0.903997	1.000000

## Feature importance pitfalls: Prediction functions

```
In [26]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.model_selection import GridSearchCV

train_X, test_X, train_y, test_y = train_test_split(X,y,test_size=0.2,
                                                    random_state=1345134)
reg = GradientBoostingRegressor(n_estimators=200,subsample=0.5,
                                random_state=3141)
grid = GridSearchCV(reg,
                    param_grid={
                        'max_leaf_nodes':[10,25,50],
                        'min_samples_leaf':[10,25,50]
                    }, n_jobs=5, cv=5)
grid = grid.fit(train_X, train_y)
```

```
In [27]: imp = pd.DataFrame({'Feature': ['$x_{}'.format(i) for i in range(1,5)],
                             'GBM Importance':grid.best_estimator_.feature_importances_})
```

## Feature Importance: Permutation Importance implementation

```
In [28]: import eli5
from eli5.sklearn import PermutationImportance

perm = PermutationImportance(grid)
perm.fit(test_X, test_y)

imp['GBM Permutation'] = perm.feature_importances_
```

## Feature Importance: Elastic Net

```
In [29]: from sklearn.linear_model import ElasticNetCV

enet = ElasticNetCV(l1_ratio=np.array([.1, .5, .7, .9, .95, .99, 1]),
                    n_alphas=500,
                    normalize=True,
                    selection='random',
                    cv=5,
                    random_state=1351)
enet.fit(train_X, train_y)

imp['Elastic Net $\\|w\\|$'] = np.abs(enet.coef_)

perm2 = PermutationImportance(enet)
perm2.fit(test_X, test_y)

imp['Elastic Net Permutation'] = perm2.feature_importances_
```

## Feature Importance: Results

- Recall  $E[y \mid x_1, z] = -\frac{1}{2}x_1 + z$  However, we don't observe  $z$  -- instead have 3 indepedent noisy observations  $x_2, x_3, x_4$ . Trying to estimate  $E[y \mid x_1, \cdots, x_4]$

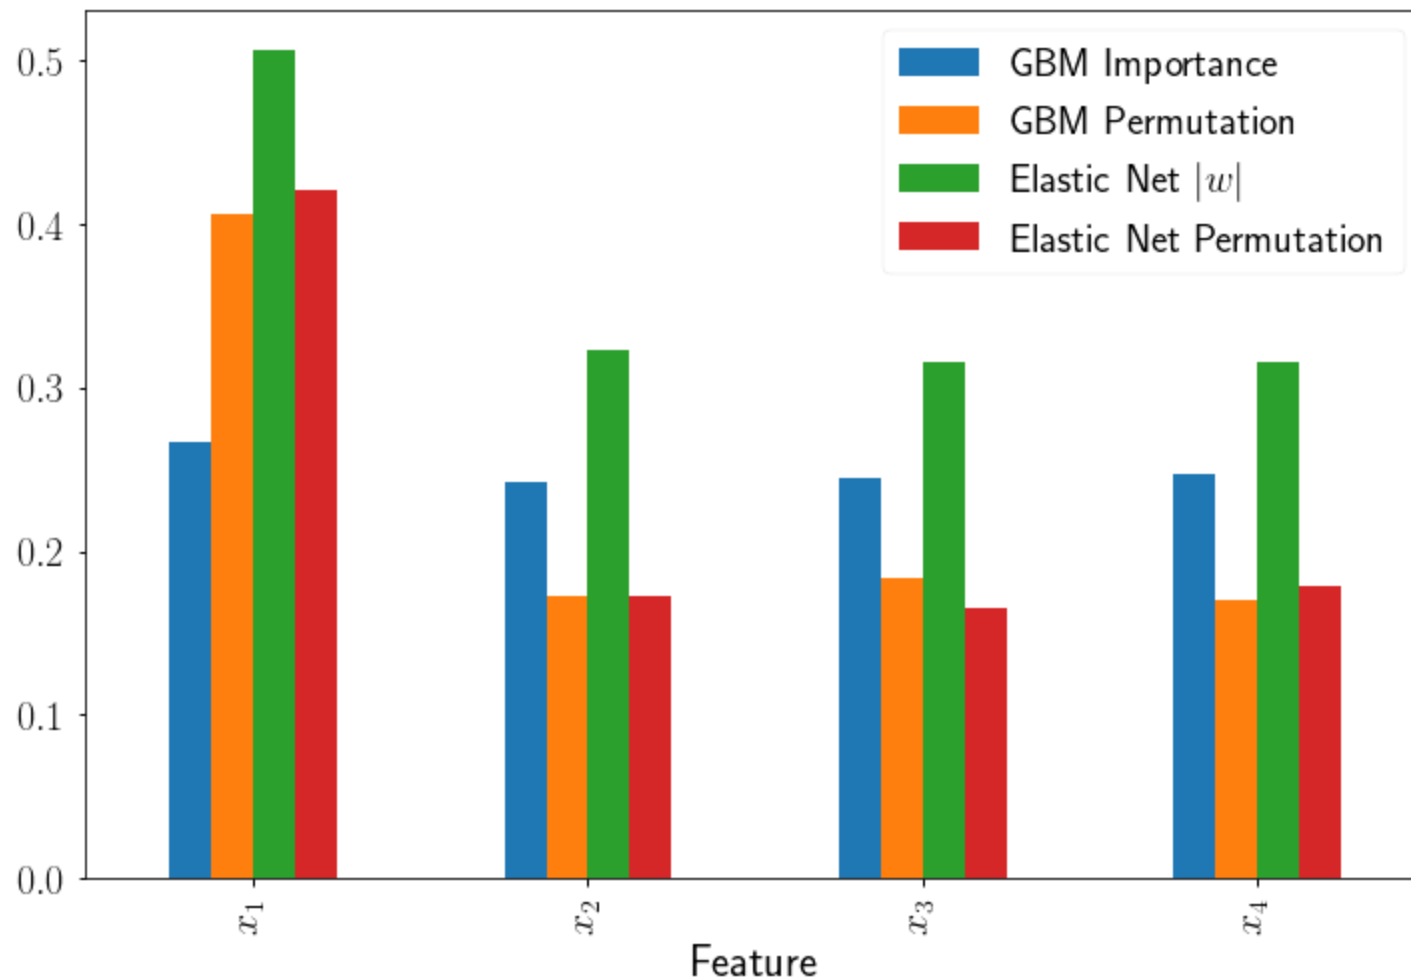
```
In [30]: hide_code_in_slideshow()
display.display(imp.set_index('Feature'))
```

	GBM Importance	GBM Permutation	Elastic Net $ w $	Elastic Net Permutation
Feature				
$x_1$	0.266826	0.406382	0.505771	0.420581
$x_2$	0.241997	0.172980	0.323082	0.172499
$x_3$	0.244454	0.183782	0.316040	0.165906
$x_4$	0.246723	0.170095	0.315358	0.178835

## Feature Importance: Results

- Recall  $E[y \mid x_1, z] = -\frac{1}{2}x_1 + z$  However, we don't observe  $z$  -- instead have 3 independent noisy observations  $x_2, x_3, x_4$ . Trying to estimate  $E[y \mid x_1, \dots, x_4]$

```
In [31]: hide_code_in_slideshow()  
sns.set_context('notebook', font_scale=2)  
fig, ax = plt.subplots(figsize=(12,8))  
plot = imp.set_index('Feature').plot(kind='bar', ax=ax)
```



- What are the implications, and what can we learn from this experiment?

## Partial Dependence Plots

- Imagine we have a subset of features we think are important.
  - This could be based on "feature importance" scores or coefficients, or other reasons (i.e. prior knowledge/research questions/etc.).
  - We may want to dig deeper to explain the relationship between our predictions and these features.
  - **Question:** Why? If we have MDI feature importance scores in hand, what do we not know?
- Directionality:
  - Note in reality we probably have complex, non-monotonic relationships
  - However even if we have a monotonic relationship, or even linear, our MDI wouldn't give any indication on the direction -- obviously weights in a linear model would.
- Partial dependence plots let us dig deeper and visualize the dependence of our prediction function (i.e. predict/predict\_proba) on one or two of these features.

## Partial Dependence Plots: Setup

- Consider an arbitrary prediction function  $f$  learned over a training set  $\mathcal{D}$ .
- This dataset includes  $N$  observations  $y_i$  of a target  $y$  for  $i = 1, 2, \dots, N$ , along with  $p$  features denoted  $x_{i,j}$  for  $j = 1, 2, \dots, p$  and  $i = 1, 2, \dots, N$ . This model generates predictions of the form:

$$\hat{y}_i = f(x_{i,1}, x_{i,2}, \dots, x_{i,p})$$

- In the case of a single feature  $x_j$ , Friedman's partial dependence plots are obtained by computing the following average and plotting it over a useful range of  $x$  values:

$$\phi_j(x) = \frac{1}{N} \sum_{i=1}^N f(x_{i,1}, \dots, x_{i,j-1}, x, x_{i,j+1}, \dots, x_{i,p})$$

- PDPs for more than one variable are computed (and plotted) similarly.
- **Question:** What would  $\phi_j(x)$  be for a linear prediction function?

## Partial Dependence Plots: Setup

$$\phi_j(x) = \frac{1}{N} \sum_{i=1}^N f(x_{i,1}, \dots, x_{i,j-1}, x, x_{i,j+1}, \dots, x_{i,p})$$

- **Question:** Think about  $\phi_j(x)$  -- what issues can we already anticipate?
- Imagine  $P(x_j = x_{j+1}) = 1$ . Then even though there is zero density where  $x_j \neq x_{j+1}$ , we still can evaluate  $\phi_j(x)$  at  $x \neq x_{i,j+1}$  for all  $i \in 1, \dots, n$
- **Previewing what's next:** Think about the potential impact of taking the average  $\frac{1}{N} \sum_{i=1}^N (\cdot)$

## PDP Example: California Housing

- Example derived from sklearn PDP example, using sklearn utilities.
- While not shown, we first train a GBM on the California Housing dataset:

```
In [33]: import numpy as np
import matplotlib.pyplot as plt

from mpl_toolkits.mplot3d import Axes3D

from sklearn.model_selection import train_test_split
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.ensemble.partial_dependence import plot_partial_dependence
from sklearn.ensemble.partial_dependence import partial_dependence
from sklearn.datasets.california_housing import fetch_california_housing

cal_housing = fetch_california_housing()

# split 80/20 train-test
X_train, X_test, y_train, y_test = train_test_split(cal_housing.data,
                                                    cal_housing.target,
                                                    test_size=0.2,
                                                    random_state=1)

names = cal_housing.feature_names

reg = GradientBoostingRegressor(n_estimators=200, subsample=0.5,
                               random_state=3141)
reg = reg.fit(X_train, y_train)
```

```
In [34]: hide_code_in_slideshow()  
print(cal_housing.DESCR)
```

California housing dataset.

The original database is available from StatLib

<http://lib.stat.cmu.edu/datasets/>

The data contains 20,640 observations on 9 variables.

This dataset contains the average house value as target variable and the following input variables (features): average income, housing average age, average rooms, average bedrooms, population, average occupation, latitude, and longitude in that order.

References

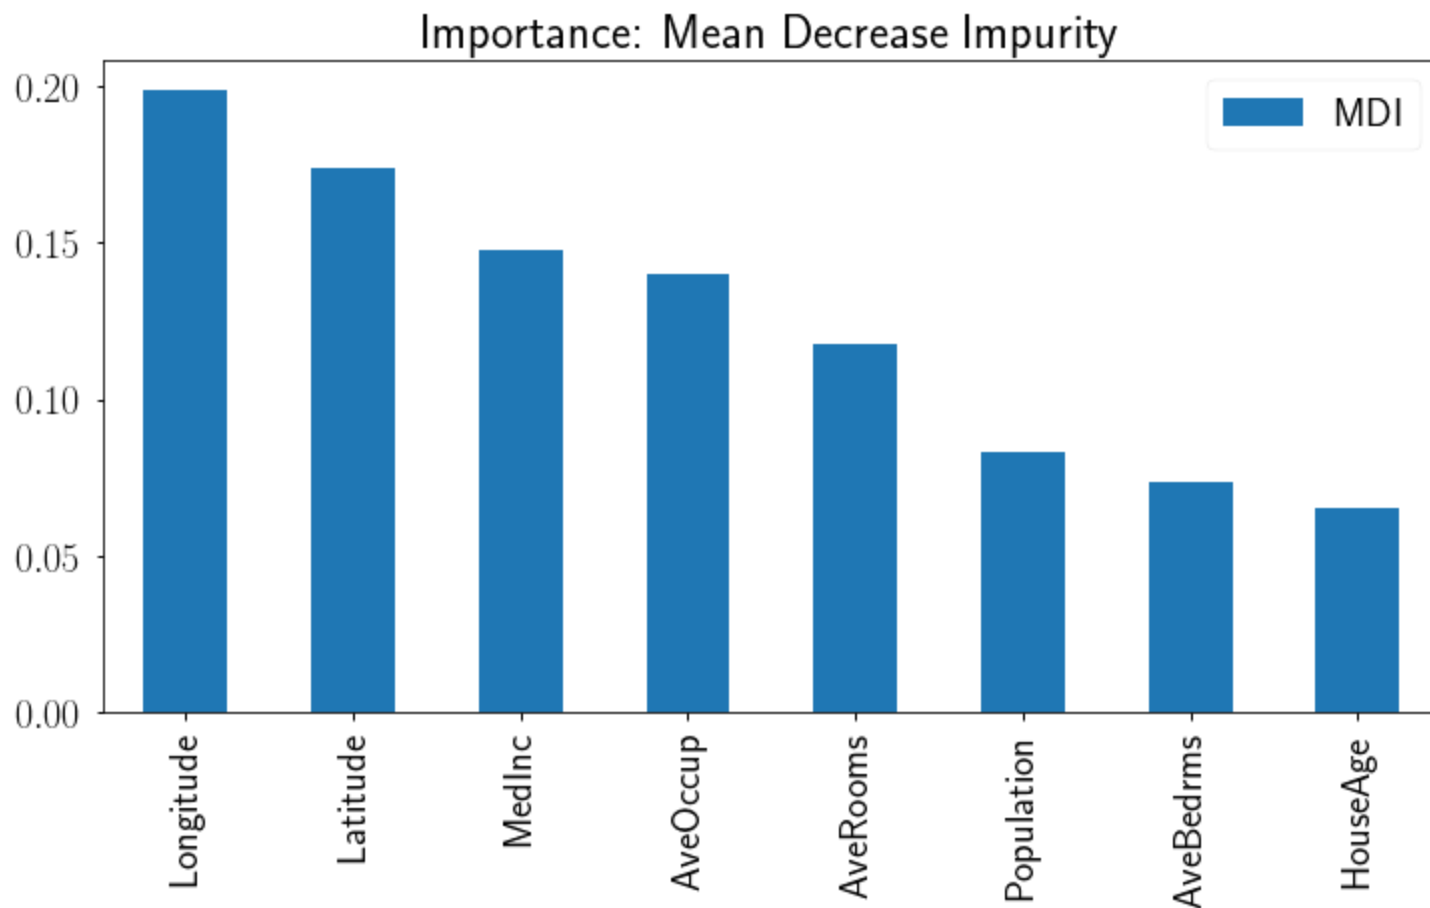
-----

Pace, R. Kelley and Ronald Barry, Sparse Spatial Autoregressions, Statistics and Probability Letters, 33 (1997) 291-297.

## PDP Example: California Housing

- Example derived from sklearn PDP example, using sklearn utilities.
- While not shown, we first train a GBM on the California Housing dataset:

```
In [35]: hide_code_in_slideshow()
fig, ax = plt.subplots(figsize=(12,6))
plot = pd.Series(reg.feature_importances_,
                  index=names).rename('MDI').to_frame()\
                  .sort_values('MDI', ascending=False).plot(kind='bar',ax=ax,
                                                             title='Importance: Mean Decrease Impurity')
```



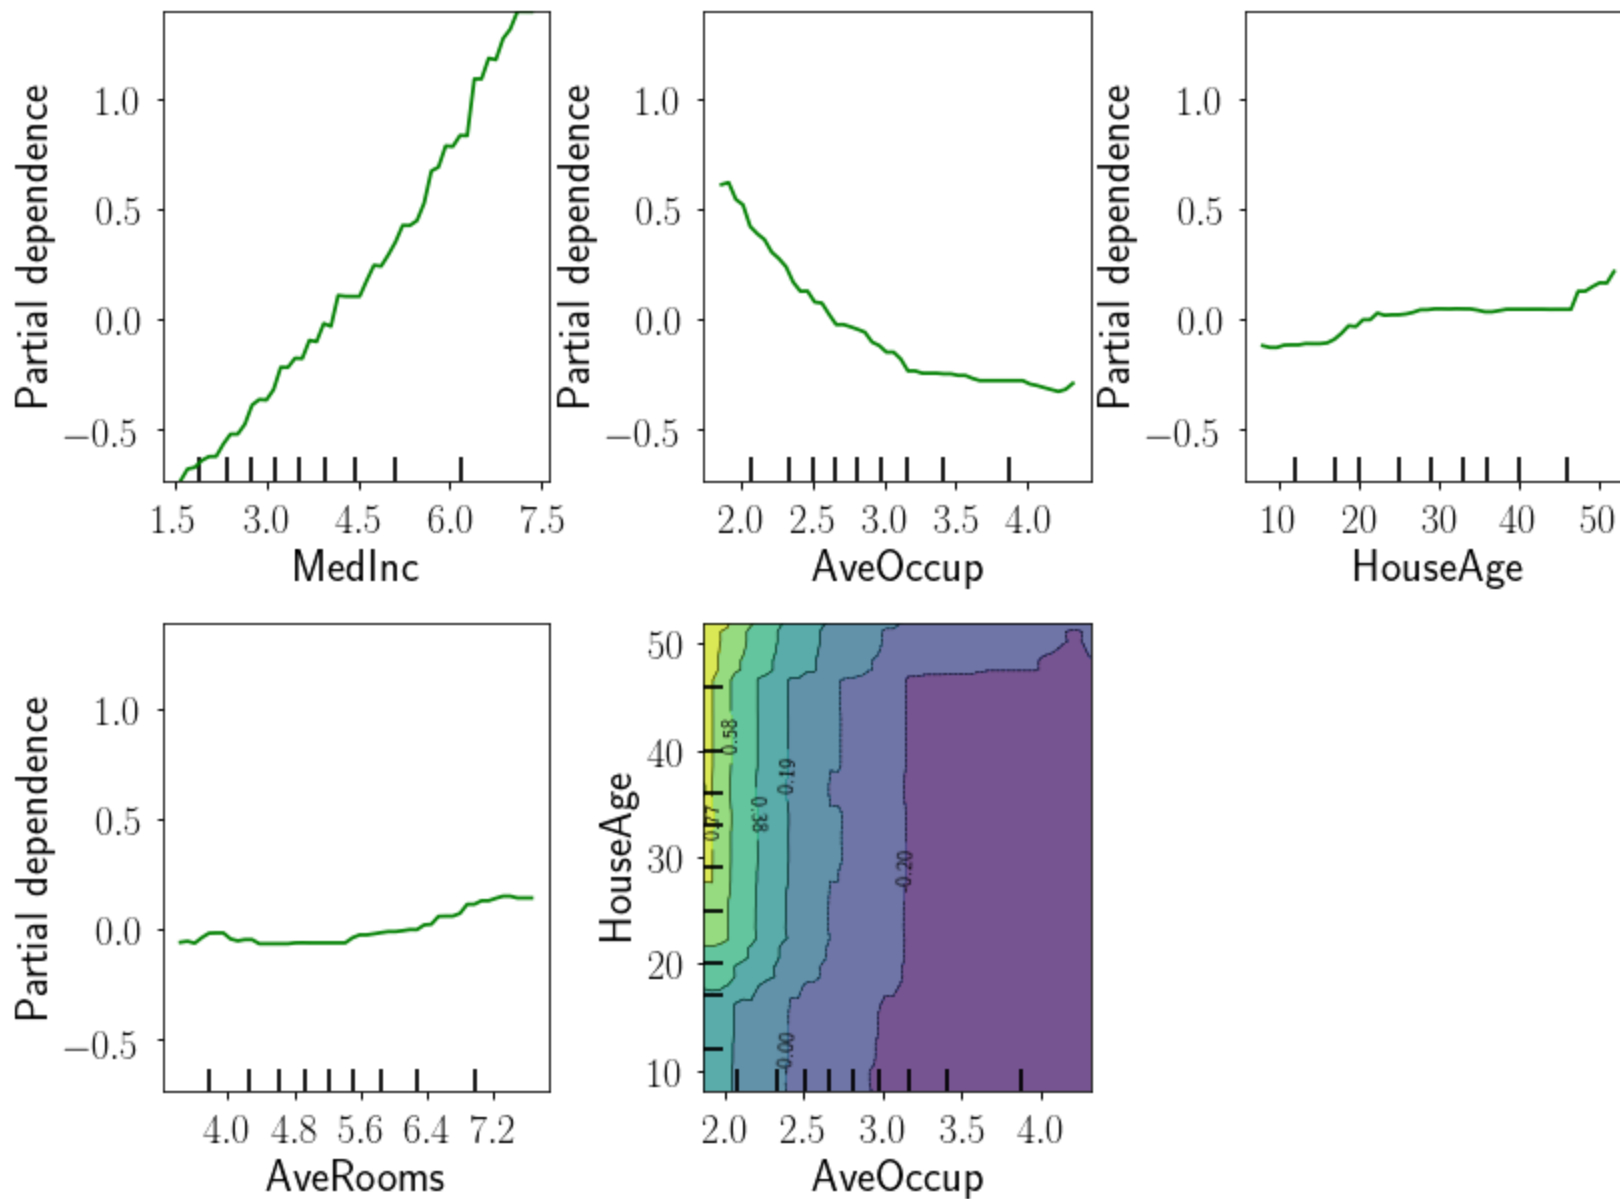
**PDP Example: California Housing**

- Next, we present the PDPs for several single features, and one pair of features.



```
In [36]: hide_code_in_slideshow()
features = [0, 5, 1, 2, (5, 1)]
fig, ax = plt.subplots(figsize=(12,10))
fig, ax = plot_partial_dependence(reg, X_train, features,
                                feature_names=names,
                                n_jobs=3, grid_resolution=50,
                                ax=ax)
fig.suptitle('Partial dependence of house value on nonlocation features')
plt.subplots_adjust(top=0.9)  # tight_layout causes overlap with suptitle
plt.show()
```

## Partial dependence of house value on nonlocation features



## PDP Example: California Housing

- Note we don't plot the spatial features -- in fact, what would perhaps be most interesting (and a potential optional extension for individual exploration) would be to plot the long/lat PDP over a base map.
  - An easy python first pass might simply use geopandas .

## PDP Pitfalls:

- Consider the following PDP for a regression problem with  $\mathcal{Y} = \mathbb{R}$  and  $X = \mathbb{R}^3$ .

```
In [37]: size = 10000
X1 = np.random.uniform(-1,1,size)
X2 = np.random.uniform(-1,1,size)
X3 = np.random.uniform(-1,1,size)
eps = np.random.normal(0,0.5,size)

Y = 0.2*X1 - 5*X2 + 10*X2*np.where(X3>=0,1,0) + eps


Xs = np.stack([X1,X2,X3]).T
```

```
In [38]: clf = GradientBoostingRegressor(max_depth=3)
params = {'n_estimators':[50,100,200]}
grid = GridSearchCV(clf, param_grid = params, n_jobs = 5, cv = 5)
grid = grid.fit(Xs,Y)
```

```
In [39]: hide_code_in_slideshow()
fig, ax = plt.subplots(figsize=(12,8))
plot_partial_dependence(grid.best_estimator_, Xs, [1],
                        n_jobs=1, grid_resolution=25, ax=ax)

plt.xlim(-1,1)
plt.ylim(-6,6)
plt.xlabel('$X_1$')
```

**pass**



- **Question:** What can we learn about the relationship between feature  $X_2$  and our target  $Y$ ?


## PDP Pitfalls:

- Here was the data generating process:

```
In [40]: size = 10000
X1 = np.random.uniform(-1,1,size)
X2 = np.random.uniform(-1,1,size)
X3 = np.random.uniform(-1,1,size)
eps = np.random.normal(0,0.5,size)
Y = 0.2*X1 - 5*X2 + 10*X2*np.where(X3>=0,1,0) + eps

Xs = np.stack([X1,X2,X3]).T
```

```
In [41]: hide_code_in_slideshow()
fig, ax = plt.subplots(figsize=(12,8))
plt.scatter(X2[X3>=0], Y[X3>=0], color='blue', label='$X_3 \ge 0$')
plt.scatter(X2[X3<0], Y[X3<0], color='green', label='$X_3 < 0$')
plt.ylabel('$Y$')
plt.xlabel('$X_2$')
plt.legend()
plt.show()
```



## PDP Pitfalls:

- Recall

$$y = 0.2x_1 - 5x_2 + 10x_2 \mathbb{1}(x_3 \geq 0) + \varepsilon$$

- **Question:** Why does the PDP fail to reveal the dependence of  $y$  on  $x_2$ ?

In his gradient boosting paper, Friedman identified this issue (here presented in our notation):

- In general, the functional form of  $\phi_j(x)$  will depend on the particular values chosen for  $\setminus j$ . If, however, this dependence is **not to strong** than the average function can represent a useful summary of the partial dependence of  $f$  on the chosen variable subset  $j$ .

## PDP Pitfalls: Another option

- One response to this shortcoming has been development of "Individual Conditional Expectation" plots -- see [Goldstein et al, 2014](https://arxiv.org/pdf/1309.6392.pdf) (<https://arxiv.org/pdf/1309.6392.pdf>).
- Their insight was that taking an average is a post-processing step that can mask structure.
- We'll use the [SauceCat/PDPbox](https://github.com/SauceCat/PDPbox) (<https://github.com/SauceCat/PDPbox>) implementation.

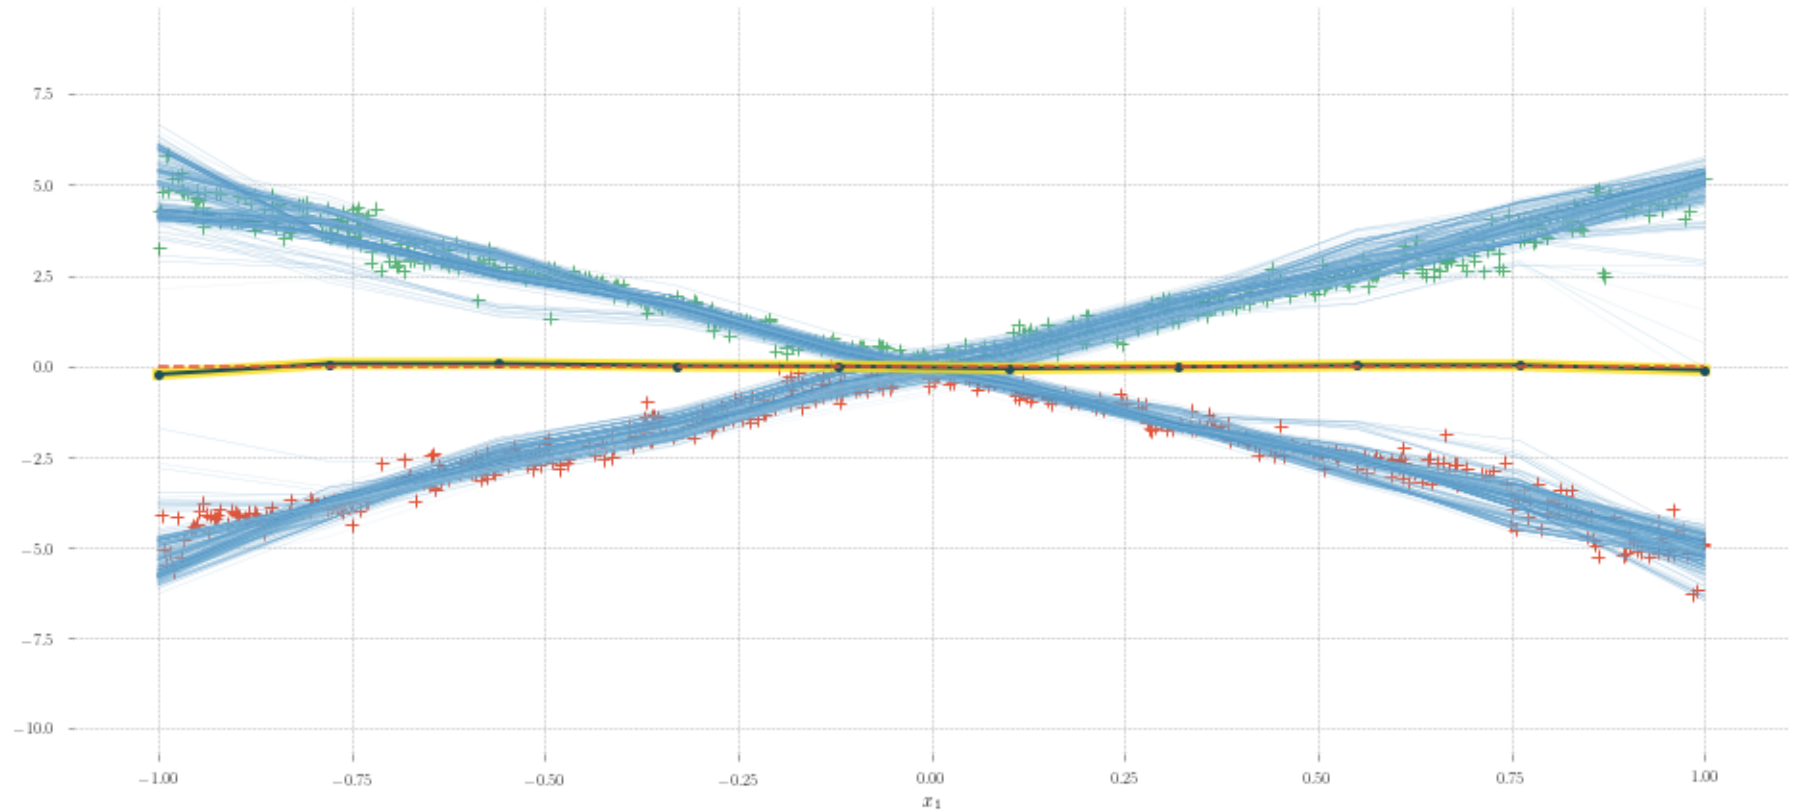
```
In [42]: hide_code_in_slideshow()
from pdpbox import pdp

train = pd.DataFrame(Xs)
train.columns = [str(x) for x in train.columns]

pdp_one = pdp.pdp_isolate(grid, train, '1')
pdp.pdp_plot(pdp_one, '$x_1$',
              plot_ort_pts=True, plot_lines=True,
              center=False, frac_to_plot=1000)
```

### ICEplot for $x_1$

Number of unique grid points: 10



- **Question:** Note the average (i.e. classic PDP plot) is shown in yellow. How do the additional traces aide clarify the dependence of  $y$  on  $x_1$ ?



## Next steps

- We started by talking about feature importance.
  - This gave a set of methods for trying to build insight into how much each feature contributed to the prediction function -- *though we saw this is somewhat ill-defined*. It didn't give insight into even the directionality of simple linear relationships -- let alone more complex structure.
- Next we looked at PDPs
  - PDPs give some insight into average relationship between one or two predictors and the target. However, this was shown to be potentially misleading, if the average relationship masks interactions.
- This led us to ICEplots, which plot a trace for each data instance separately.
- Note this trajectory (moving from global to data-instance-specific measures of importance) could be extended -- next steps might look at additive feature attribution methods such as SHAP values (<https://arxiv.org/pdf/1705.07874.pdf>).

## Bibliography

### Papers

- Brieman, L (2001). *Random Forests*. (<https://link.springer.com/content/pdf/10.1023/A:1010933404324.pdf>)
- Friedman, J (2001). *Greedy Function Approximation: A Gradient Boosting Machine*. (<https://statweb.stanford.edu/~jhf/ftp/trebst.pdf>)
- Goldstein et al (2014). *Peeking Inside the Black Box: Visualizing Statistical Learning with Plots of Individual Conditional Expectation* (<https://arxiv.org/pdf/1309.6392.pdf>)

### Packages

- `scikit-learn/scikit-learn` (<https://github.com/scikit-learn/scikit-learn>)
- `TeamHG-Memex/eli5` (<https://github.com/TeamHG-Memex/eli5>)
- `SauceCat/PDPbox` (<https://github.com/SauceCat/PDPbox>)