

# Information Theory

David Rosenberg

New York University

November 1, 2015

# A Measure of Information?

- Consider a discrete random variable  $X$ .
- How much “information” do we gain from observing  $X$ ?
- Information  $\approx$  “degree of surprise” from observing  $X = x$ .
- If we know  $\mathbb{P}(X = 0) = 1$ , then observing  $X = 0$  gives no information.
- If we know  $\mathbb{P}(X = 0) = .999$ :
  - Observing  $X = 0$  gives little information.
  - Observing  $X = 1$  gives a lot of surprise / “information”
- Information measure  $h(x)$  should depend on  $p(x)$ :
  - Smaller  $p(x) \implies$  More information  $\implies$  Larger  $h(x)$

# Shannon Information Content of an Outcome

## Definition

Let  $X \in \mathcal{X}$  have PMF  $p(x)$ . The **Shannon information content** of an **outcome**  $x$  is

$$h(x) = \log \left( \frac{1}{p(x)} \right),$$

where the base of the log is 2. Information is measured in **bits**. (Or **nats** if the base of the log is  $e$ .)

- Less likely outcome gives more information.
- Information is **additive** for independent events:
  - If  $X$  and  $Y$  are independent,

$$\begin{aligned} h(x, y) &= -\log p(x, y) = -\log [p(x)p(y)] \\ &= -\log p(x) - \log p(y) \\ &= h(x) + h(y) \end{aligned}$$

# Entropy

## Definition

Let  $X \in \mathcal{X}$  have PMF  $p(x)$ . The **entropy** of  $X$  is

$$\begin{aligned} H(X) &= \mathbb{E}_p \log \left( \frac{1}{p(X)} \right) \\ &= - \sum_{x \in \mathcal{X}} p(x) \log p(x), \end{aligned}$$

using convention that  $0 \log 0 = 0$ , since  $\lim_{x \rightarrow 0^+} x \log x = 0$ .

- Entropy of  $X$  is the expected information gain from observing  $X$ .
- Entropy only depends on distribution  $p$ , so we can write  $H(p)$ .

# Coding

## Definition

A **binary source code**  $C$  is a mapping from  $\mathcal{X}$  to finite 0/1 sequences.

- Consider r.v.  $X \in \mathcal{X}$  and binary source code  $C$  defined as:

$x$	$p(x)$	$C(x)$
1	$1/2$	0
2	$1/4$	10
3	$1/8$	110
4	$1/8$	111

# Expected Code Length

- Consider r.v.  $X \in \mathcal{X}$  and binary source code  $C$  defined as:

$x$	$p(x)$	$C(x)$	$\log \frac{1}{p(x)}$
1	1/2	0	$\log_2 2 = 1$
2	1/4	10	$\log_2 4 = 2$
3	1/8	110	$\log_2 8 = 3$
4	1/8	111	$\log_2 8 = 3$

- The **entropy** is  $H(X) = \mathbb{E} \log [1/p(x)]$ :


$$H(X) = \frac{1}{2}(1) + \frac{1}{4}(2) + \frac{1}{8}(3) + \frac{1}{8}(3) = 1.75 \text{ bits.}$$

- The **expected code length** is

$$L(C) = \frac{1}{2}(1) + \frac{1}{4}(2) + \frac{1}{8}(3) + \frac{1}{8}(3) = 1.75 \text{ bits.}$$

# Prefix Codes

- A code is a **prefix code** if no codeword is a prefix of another.
- Prefix codes can be represented on trees:



- Each leaf node is a codeword.
- It's encoding represents the path from root to leaf.

From David MacKay's *Information Theory, Inference, and Learning Algorithms*, Section 5.1.

# Data Compression: What's the Best Prefix Code?

- For  $X \sim p(x)$ , we get best compression with codeword lengths

$$\ell^*(x) \approx -\log p(x).$$

- **Optimal bit length of  $x$  is the Shannon Information of  $x$ .**
- Then the **expected codeword length** is

$$\begin{aligned} L^* &= \mathbb{E}[-\log p(X)] \\ &= H(X) \end{aligned}$$

- Entropy  $H(X)$  gives a **lower bound** on coding performance.
- Shannon's Theorem says we can achieve  $H(X)$  within 1 bit.



# Shannon's Source Coding Theorem

## Theorem (Shannon's Source Coding Theorem)

*The expected length  $L$  of any binary prefix code for r.v.  $X$  is at least  $H(X)$ :*

$$L \geq H(X).$$

*There exist codes with lengths  $\ell(x) = \lceil -\log_2 p(x) \rceil$  achieving*

$$H(X) \leq L < H(X) + 1.$$

- **Notation**  $\lceil x \rceil = \text{ceil}(x) = (\text{smallest integer } \geq x)$

# Shannon's Source Coding Theorem: Summary

- For any  $X \sim p(x)$ ,  $\exists$  code with  $L \approx H(X)$ .
- Get arbitrarily close to  $H(X)$  by grouping multiple  $X$ 's and coding all at once.
- If we know the distribution of  $X$ , we can code optimally.
  - e.g. Use **Huffman codes** or **arithmetic codes**.
- What if we don't know  $p(x)$ , and we use  $q(x)$  instead?

# Coding with the Wrong Distribution: Core Calculation

- Allow fractional code lengths:  $\ell_q(x) = -\log q(x)$
- Then expected length for coding  $X \sim p(x)$  using  $\ell_q(x)$  is

$$\begin{aligned}
 L &= \mathbb{E}_{X \sim p(x)} \ell_q(X) \\
 &= -\sum_x p(x) \log q(x) \\
 &= \sum_x p(x) \log \left[ \frac{p(x)}{q(x)} \frac{1}{p(x)} \right] \\
 &= \sum_x p(x) \log \frac{p(x)}{q(x)} + \sum_x p(x) \log \frac{1}{p(x)} \\
 &= \text{KL}(p \| q) + H(p),
 \end{aligned}$$

where  $\text{KL}(p \| q)$  is the Kullback-Leibler divergence between  $p$  and  $q$ .

# Entropy, Cross-Entropy, and KL-Divergence

- The **Kullback-Leibler** or “**KL**” **Divergence** is defined by

$$\text{KL}(p\|q) = \mathbb{E}_p \log \left( \frac{p(X)}{q(X)} \right).$$

- $\text{KL}(p\|q)$ : **#(extra bits)** needed if we code with  $q(x)$  instead of  $p(x)$ .
- The **cross entropy** for  $p(x)$  and  $q(x)$  is defined as

$$H(p, q) = -\mathbb{E}_p \log q(X).$$

- $H(p, q)$ : **#(bits)** needed to code  $X \sim p(x)$  using  $q(x)$ .
- Summary:

$$H(p, q) = H(p) + \text{KL}(p\|q).$$

# Coding with the Wrong Distribution: Integer Lengths

## Theorem

*If we code  $X \sim p(x)$  using code lengths  $\ell(x) = \lceil -\log_2 q(x) \rceil$ , the expected code length is bounded as*

$$H(p) + KL(p\|q) \leq \mathbb{E}_p \ell(X) < H(p) + KL(p\|q) + 1.$$

- So with an implementable code (using integer codeword lengths), the expected code length is within 1 bit of what could be achieved with  $\ell(x) = -\log_2 q(x)$ .
- Proof is a slight tweak on the “core calculation”.

# Jensen's Inequality

## Theorem (Jensen's Inequality)

If  $f : \mathcal{X} \rightarrow \mathbf{R}$  is a **convex** function, and  $X \in \mathcal{X}$  is a random variable, then

$$\mathbb{E}f(X) \geq f(\mathbb{E}X).$$

Moreover, if  $f$  is **strictly convex**, then equality implies that  $X = \mathbb{E}X$  with probability 1 (i.e.  $X$  is a constant).

- e.g.  $f(x) = x^2$  is convex. So  $\mathbb{E}X^2 \geq (\mathbb{E}X)^2$ . Thus

$$\text{Var}X = \mathbb{E}X^2 - (\mathbb{E}X)^2 \geq 0.$$

# Gibbs Inequality ( $KL(p\|q) \geq 0$ )

## Theorem (Gibbs Inequality)

Let  $p(x)$  and  $q(x)$  be PMFs on  $\mathcal{X}$ . Then

$$KL(p\|q) \geq 0,$$

with equality iff  $p(x) = q(x)$  for all  $x \in \mathcal{X}$ .

- KL divergence measures the “distance” between distributions.
- Note:
  - KL divergence **not a metric**.
  - KL divergence is **not symmetric**.

# Gibbs Inequality: Proof

$$\begin{aligned}
 \text{KL}(p\|q) &= \mathbb{E}_p \left[ -\log \left( \frac{q(X)}{p(X)} \right) \right] \\
 &\geq -\log \left[ \mathbb{E}_p \left( \frac{q(X)}{p(X)} \right) \right] \quad (\text{Jensen's}) \\
 &= -\log \left[ \sum_{\{x|p(x)>0\}} p(x) \frac{q(x)}{p(x)} \right] \\
 &= -\log \left[ \sum_{x \in \mathcal{X}} q(x) \right] \\
 &= -\log 1 = 0.
 \end{aligned}$$

- Since  $-\log$  is strictly convex, we have strict equality iff  $q(x)/p(x)$  is a constant, which implies  $q = p$ .
- Essentially the same proof for PDFs.



# KL-Divergence for Model Estimation

- Suppose  $\mathcal{D} = \{x_1, \dots, x_n\}$  is a sample from **unknown**  $p(x)$  on  $\mathcal{X}$ .
- **Hypothesis space**:  $\mathcal{P}$  some set of distributions on  $\mathcal{X}$ .
- Idea: Find  $q \in \mathcal{P}$  that minimizes  $\text{KL}(p||q)$ :

$$\arg \min_{q \in \mathcal{P}} \text{KL}(p, q) = \arg \min_{q \in \mathcal{P}} \mathbb{E}_p \left[ \log \left( \frac{p(X)}{q(X)} \right) \right]$$

- Don't know  $p$ , so **replace expectation by average over  $\mathcal{D}$** :

$$\arg \min_{q \in \mathcal{P}} \left\{ \frac{1}{n} \sum_{i=1}^n \log \left[ \frac{p(x_i)}{q(x_i)} \right] \right\}$$

# Estimated KL-Divergence

- The **estimated KL-divergence**:

$$\begin{aligned} & \frac{1}{n} \sum_{i=1}^n \log \left[ \frac{p(x_i)}{q(x_i)} \right] \\ &= \frac{1}{n} \sum_{i=1}^n \log p(x_i) - \frac{1}{n} \sum_{i=1}^n \log q(x_i). \end{aligned}$$

- The minimizer of this over  $q \in \mathcal{P}$  is also

$$\arg \max_{q \in \mathcal{P}} \sum_{i=1}^n \log q(x_i).$$

- This is exactly the objective for the **MLE**.
- **Minimizing KL between model and truth leads to MLE.**