Gradient Boosting Practice: Poisson Response

David S. Rosenberg

Suppose we're trying to predict a distribution of count from some input covariates. The simplest distribution in this situation is the Poisson distribution:

$$p(k;\lambda) = \frac{e^{-\lambda}\lambda^k}{k!}$$

on $k = 0, 1, 2, 3, ... \lambda \in (0, \infty)$.

1 Linear Conditional Probability Model

- Input: $x \in \mathbf{R}^d$.
- Output: $y \in \{0, 1, 2, ...\}$
- Data:

$$\mathcal{D} = ((x_1, y_1), \dots, (x_n, y_n)) \in \left(\mathbf{R}^d \times \{0, 1, 2, \dots\}\right)^n$$

assume is sampled i.i.d. from some distribution $P_{\mathfrak{X}\times\mathfrak{Y}}$.

• Action: $\lambda \in (0, \infty)$, where λ is the parameter of a Poisson distribution.

We've got to map input x to action λ in our action space, which is $(0, \infty)$.

$$x \mapsto \underbrace{w^{\mathsf{T}} x}_{\text{score } s \in \mathbf{R}} \mapsto \underbrace{\lambda}_{\in (0, \infty)}$$

To map the score s into our action space, we could use the transfer function $\psi(s) = \exp(s)$. Then

$$\lambda = \exp\left(w^{\mathsf{T}}x\right).$$

So if we predict λ , what's the probability of an observed count k for input vector \mathbf{x} ?

$$p(y = k \mid x; w) = \frac{e^{-\lambda(x)}\lambda(x)^k}{k!}$$
$$= \frac{e^{-\exp(w^T x)} \left[\exp(w^T x)\right]^k}{k!}.$$

The conditional likelihood for particular example (x_i, y_i) , (where y_i is a count) is

$$p(y = y_i \mid x_i; w) = \frac{e^{-\exp(w^T x_i)} \left[\exp(w^T x_i)\right]^{y_i}}{y_i!}.$$

Easier to work with the log:

$$\log p(y = y_i | x_i; w) = -\exp(w^T x_i) + y_i w^T x_i - \log(y_i!)$$

What do we need to find to fit this model? w. our strategy is to use maximum log-likelihood:

$$\begin{aligned} \log L_{\mathcal{D}}(w) &= \log p(\mathcal{D}; w) \\ &= \sum_{i=1}^{n} \log p(y = y_i \mid x_i; w) \\ &= \sum_{i=1}^{n} \left[-\exp\left(w^T x_i\right) + y_i w^T x_i - \log(y_i!) \right] \end{aligned}$$

So find w maximizing this log-likelihood and we're done. Can use standard gradient based methods.

2 Nonlinear approach

In a nonlinear approach, we'll replace the linear score function $s = w^T x$ with a nonlinear function s = f(x):

$$x \mapsto \underbrace{f(x)}_{\text{score } s \in \mathbf{R}} \mapsto \underbrace{\lambda}_{\in (0,\infty)}.$$

Again, we can use the transfer function $\psi(s) = \exp(s)$. So

$$\lambda = \exp(f(x))$$
.

For score function f, the probability of $y_i \mid x_i$ is:

$$p(y = y_i \mid x_i; f) = \frac{e^{-\exp(f(x_i))} [\exp(f(x_i))]^{y_i}}{y_i!}.$$

Easier to work with the log:

$$\log p(y = y_i | x_i; f) = -\exp(f(x_i)) + y_i f(x_i) - \log(y_i!)$$

Somehow we want to find a function f that gives high log-likelihood to our observed data:

$$\log L_{\mathcal{D}}(w) = \sum_{i=1}^{n} \left[-\exp(f(x_i)) + y_i f(x_i) - \log(y_i!) \right]$$

3 Gradient Boosting Approach

Let's differentiate $\log p(y = y_i \mid x_i; f)$ w.r.t. $f(x_i)$:

$$\frac{\partial}{\partial f(x_i)} \log p(y = y_i \mid x_i; f) = -\exp(f(x_i)) + y_i$$

Now differentating the full log-likelihooed is

$$\frac{\partial}{\partial f(x_i)} [\log L_{\mathcal{D}}(f)] = \frac{\partial}{\partial f(x_i)} [-\exp(f(x_i)) + y_i f(x_i) - \log(y_i!)]$$
$$= -\exp(f(x_i)) + y_i$$

So optimal unconstrained step direction for changing the vector of evaluations $\mathbf{f} = (f(x_1), \dots, f(x_n))$ is

$$-\mathbf{g} = \left(-y_1 + \exp\left(f(x_1)\right), \dots, -y_n + \exp\left(f(x_n)\right)\right)$$

Fix some base hypothesis space \mathcal{H} of functions $h: \mathbf{R}^d \to \mathbf{R}$. Then, our actual step direction will be the $h \in \mathcal{H}$ that best fits $-\mathbf{g}$ in the least squares sense:

$$\begin{aligned} & \underset{h \in \mathcal{H}}{\text{arg min}} \sum_{i=1}^{n} \left(-\mathbf{g}_{i} - h(x_{i}) \right)^{2} \\ &= & \underset{h \in \mathcal{H}}{\text{arg min}} \sum_{i=1}^{n} \left(\left[-y_{i} + \exp\left(f(x_{i})\right) \right] - h(x_{i}) \right)^{2} \end{aligned}$$

So to recap:

- 1. Up to this point, our score function is f.
- 2. We want to improve f.
- 3. The optimal step direction for $f(x_i)$ is $-y_i + \exp(f(x_i))$. We can evaluate this. It's a real number.
- 4. So we have a bunch of $(x_i, -\mathbf{g}_i)$ pairs that we will use regression over \mathcal{H} to fit.

Then we add something like 0.1h to f and repeat.